Windows Azure

Python and Windows Azure

Wenming Ye
Sr. Program Manager
Microsoft Research

Twitter: @wenmingye

Cloud Computing

IaaS

Infrastructure-as-a-Service

host

PaaS

Platform-as-a-Service

build

SaaS

Software-as-a-Service

consume

Windows Azure

Comprehensive set of services that enable you to quickly build, deploy and manage applications across a global network of Microsoft-managed datacenters

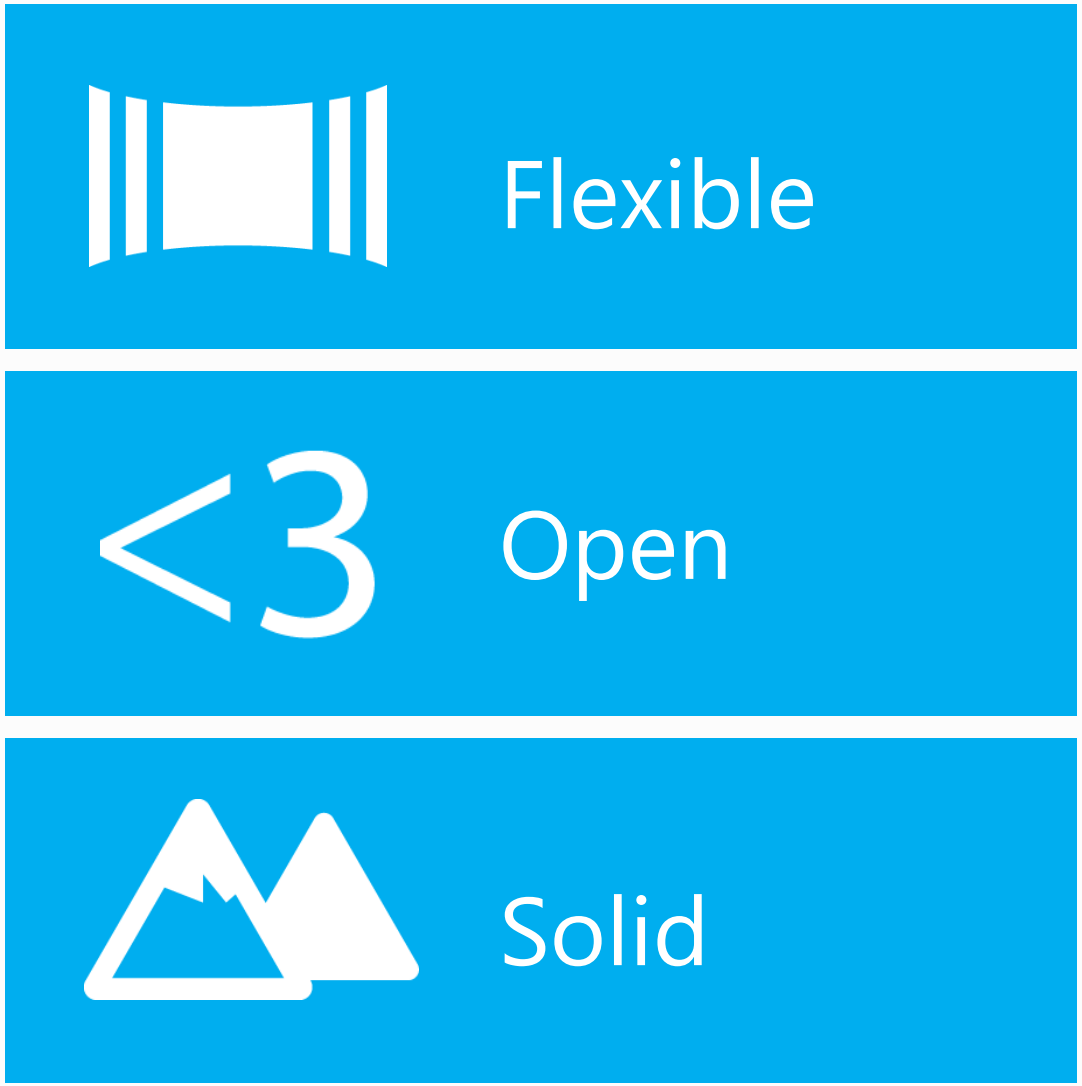

Global Footprint

Windows Azure

Virtual machines

Cloud services

Web sites

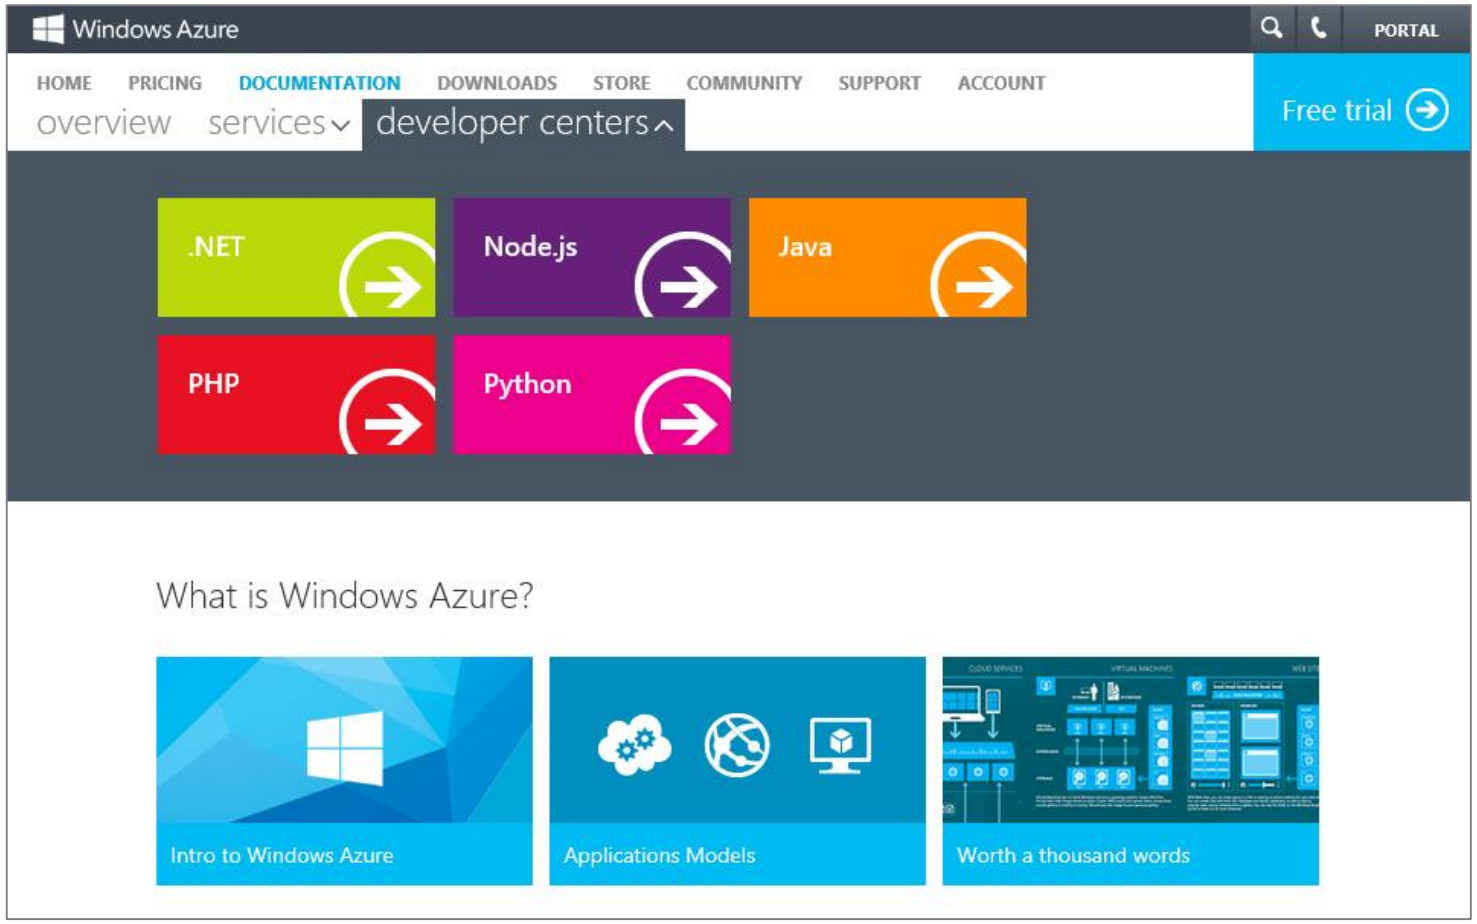

Multiple languages

→ http://WindowsAzure.com

PTVS: Python Tools for Visual Studio

Windows Azure™

Python Tools for VS:

- Project mgmt
- Solutions
- Deployment

- Built-in REPL
- IPython REPL
- Interactive Parallel Computing
- Inline graphics

- Syntax hi-lite
- Intellisense
- Find / Browse

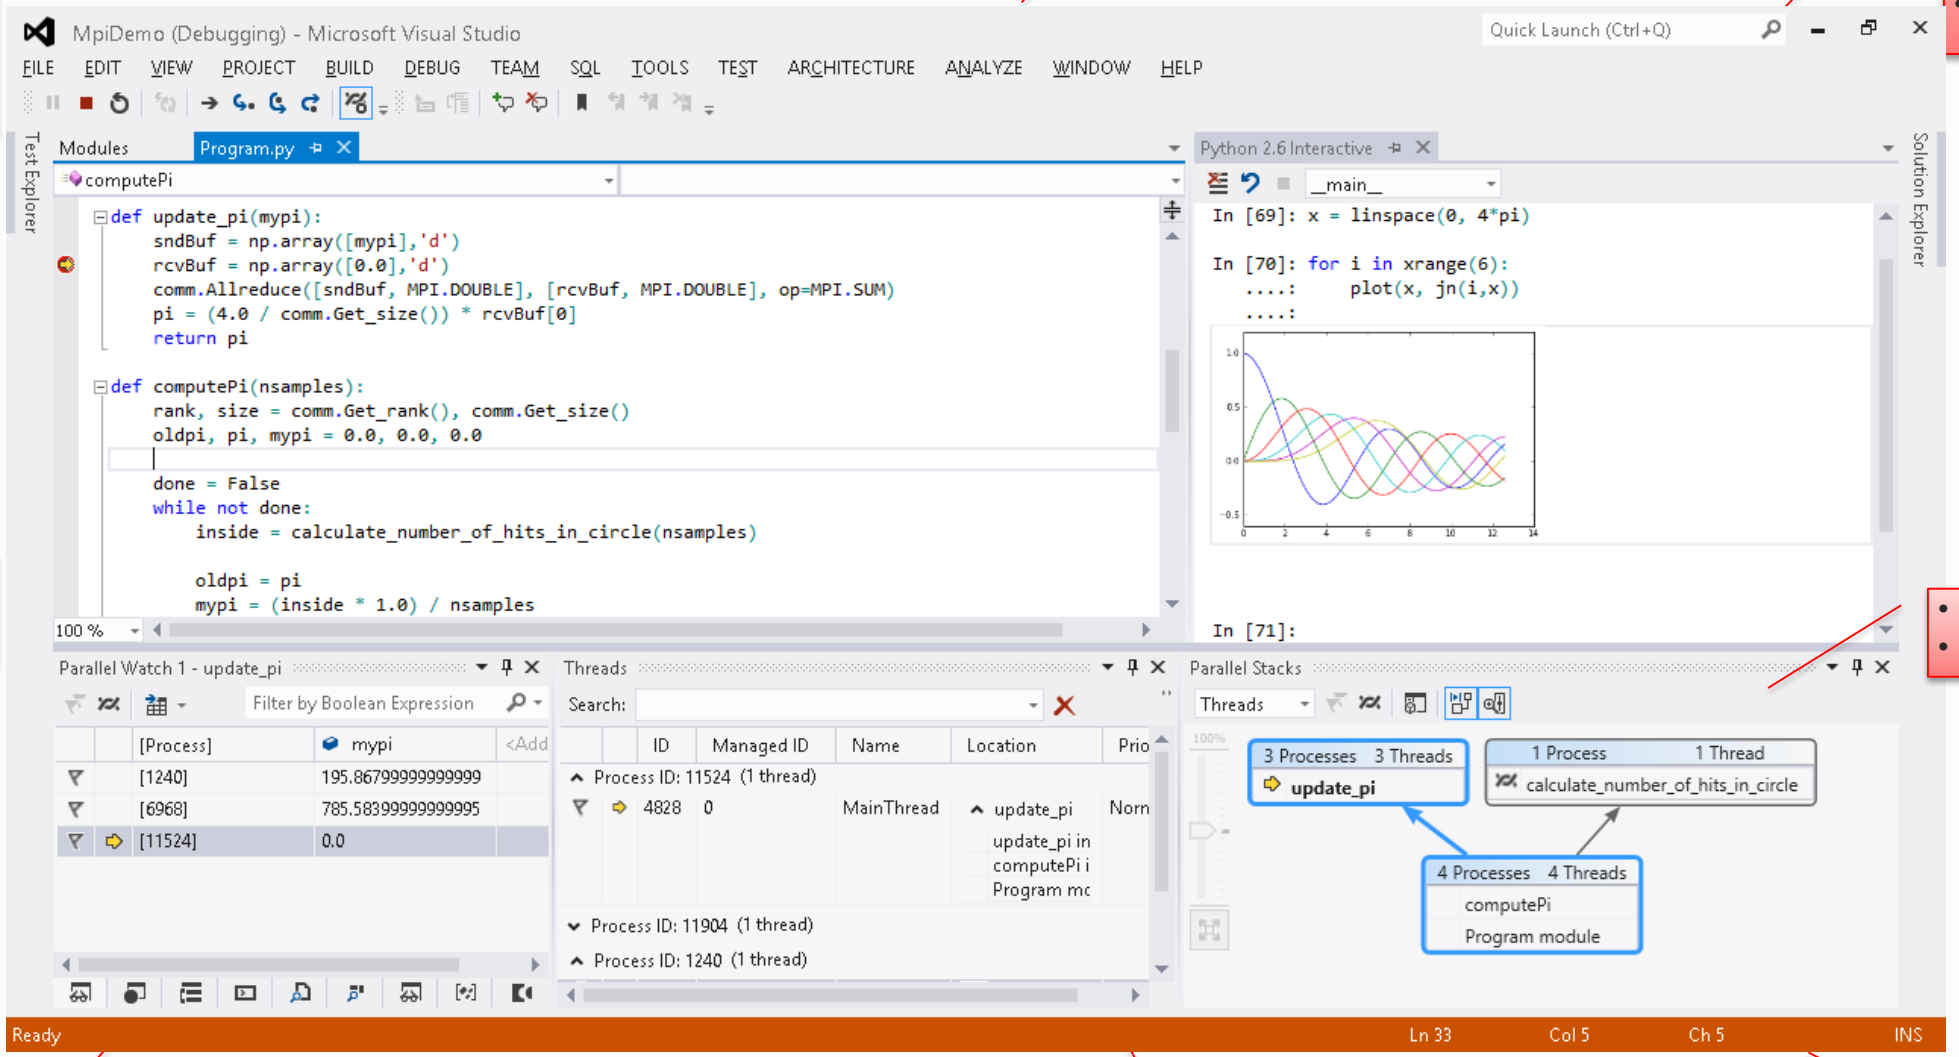

- Profiler
- A/B comparions

- Parallel Watch view
- Parallel Stack view

- CPython
- IronPython
- Or any interpreter

ws Azure

- Python Debugger
- .Net Debugger
- Remote Debugging

- HPC Support
- F5 MPI debugging
- Batch or Interactive

Typical Scenarios

- General purpose programming
  - Projects
  - Edit/Intellisense/Browse
  - Debug
  - Profile
- Cloud / Web
  - Azure Cloud Service
  - Azure Web Sites
  - Django, WFastCGI included
  - IPython notebook
- Technical / Scientific Computing
  - HPC / MPI
  - Cluster Debugging
  - Inline REPL graphics
  - IPython REPL & notebook

PTVS Features

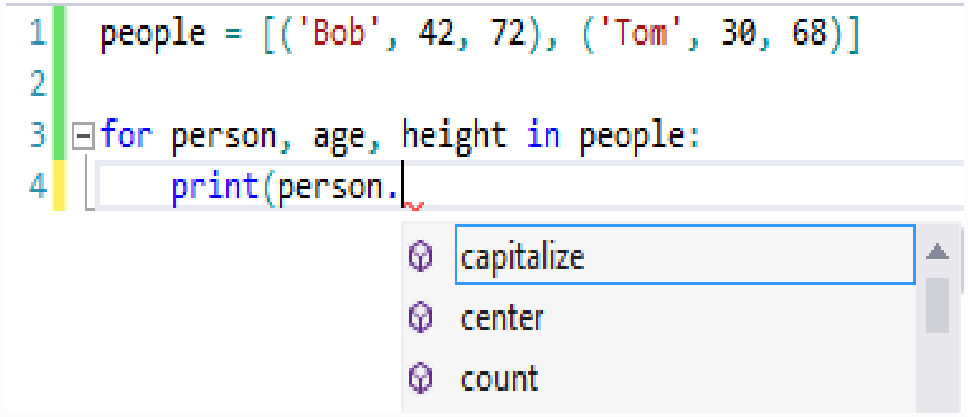

Intellisense

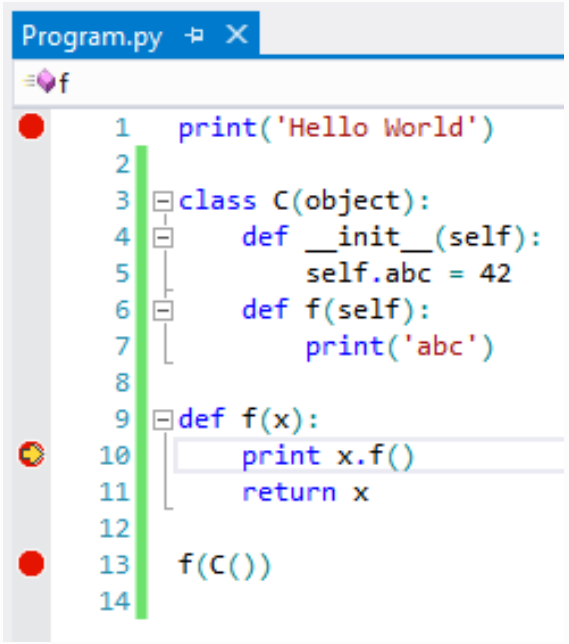

Debugging

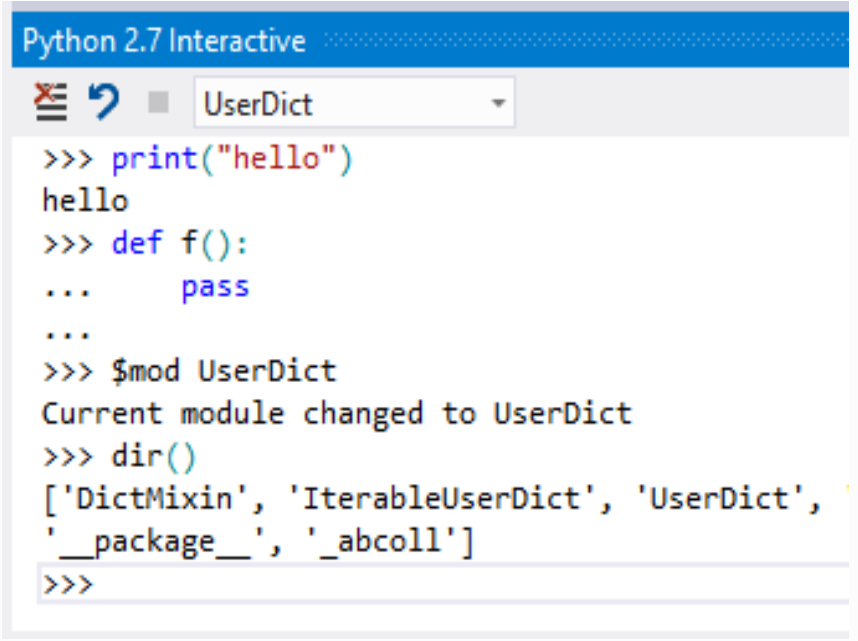

REPL – IPython too

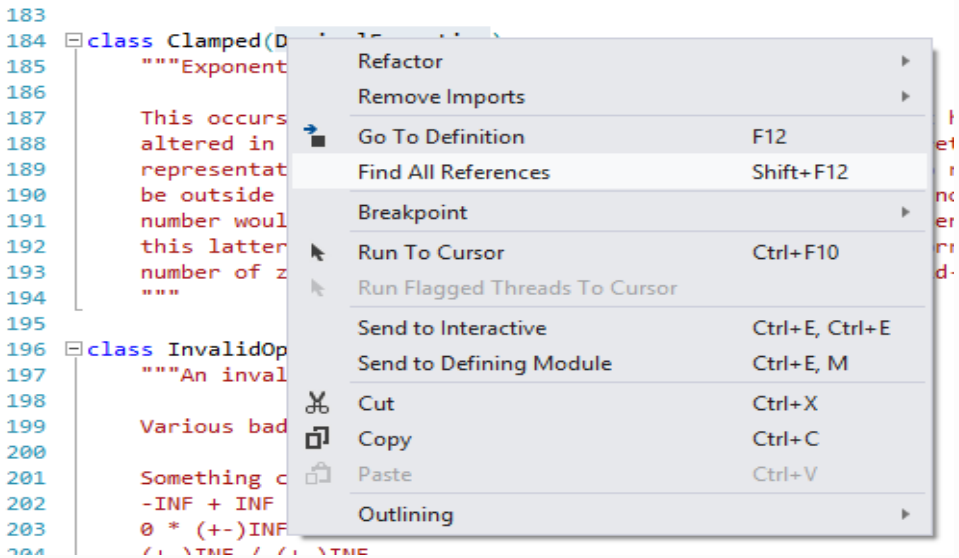

Find All Refs

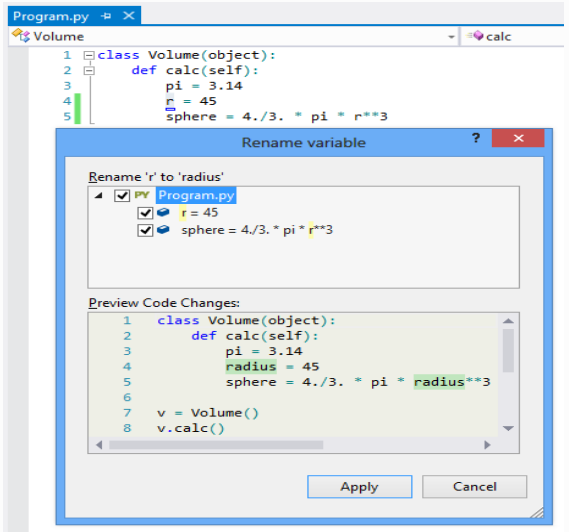

Refactor/Rename

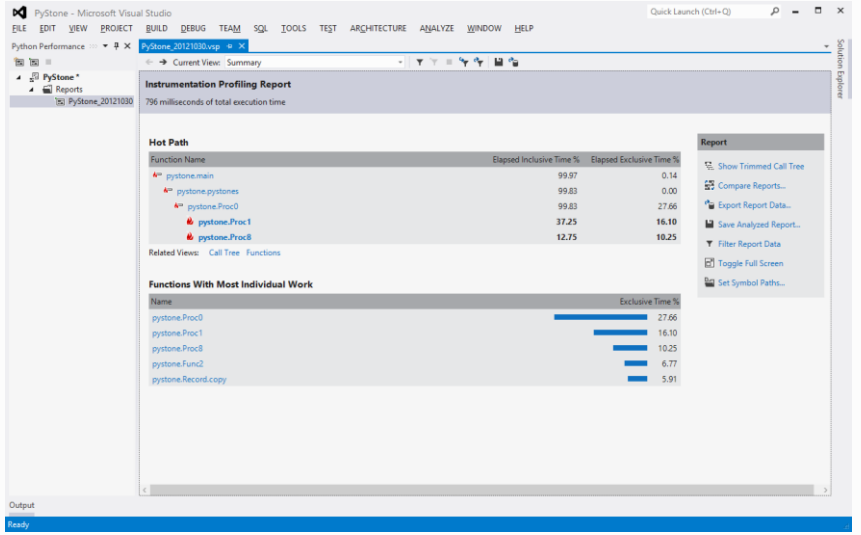

Profiling

Windows Azure

Cloud / Web

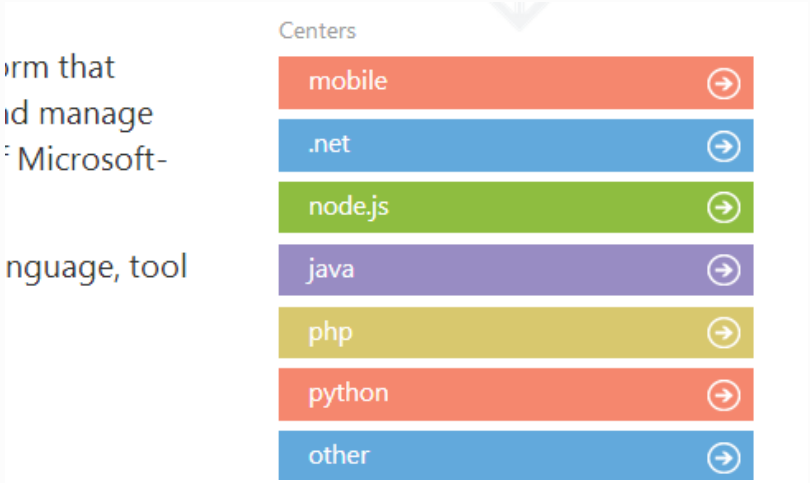

Table Storage

To ensure a table exists, call **create_table**:

```
from azure.storage import TableService
ts = TableService(account_name, account_key)
ts.create_table('tasktable')
```

A new entity can be added by calling **insert_entity**:

```
from datetime import datetime
ts = TableService(account_name, account_key)
ts.create_table('tasktable')
ts.insert_entity(
    'tasktable',
    {
       'PartitionKey' : 'tasksSeattle',
       'RowKey': '1',
       'Description': 'Take out the trash',
       'DueDate': datetime(2011, 12, 14, 12)
    }
)
```

Azure Python Client Libs for Windows, MacOS, Linux

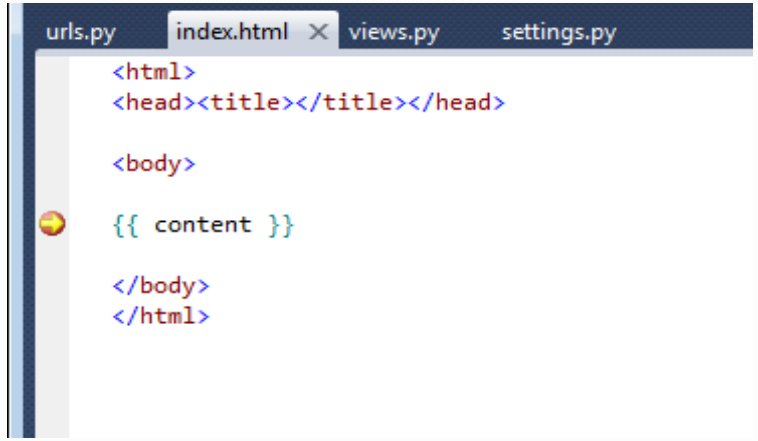

Django template debug

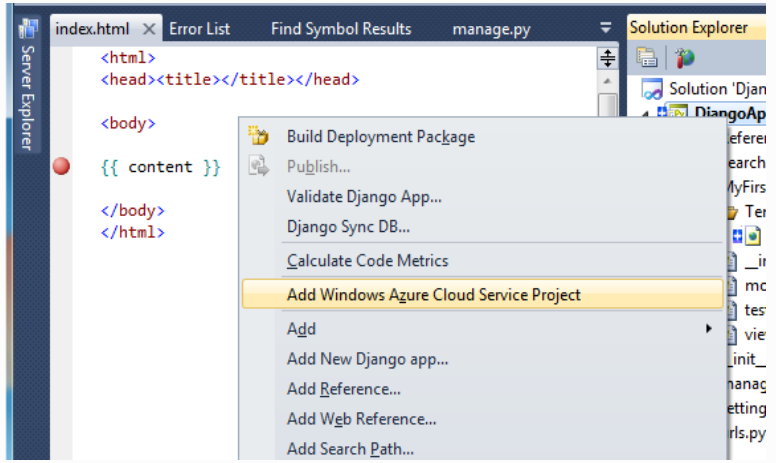

Django Deploy

2. Setup the Python Fast CGI Handler

```
%windir%\system32\inetsrv\appcmd set config
/section:system.webServer/fastCGI "/+[fullPath='c:\Python
arguments='C:\inetpub\wwwroot\wfastcgi.py']"
```

3. Register the handler for this site

```
%windir%\system32\inetsrv\appcmd set config
/section:system.webServer/handlers "/+
[name='Python_via_FastCGI',path='*',verb='*',modules='Fas
ptProcessor='c:\Python27\python.exe|C:\inetpub\wwwroot\wf
rceType='Unspecified']"
```

WFastCGI for use with Django, Flask, Bottle, Web2Py, …

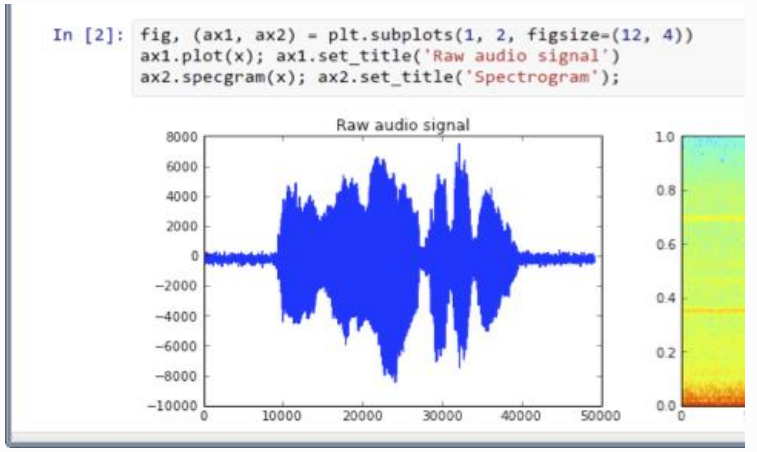

IPython Notebook: Python in the browser

Windows Azure

Technical / Scientific Computing

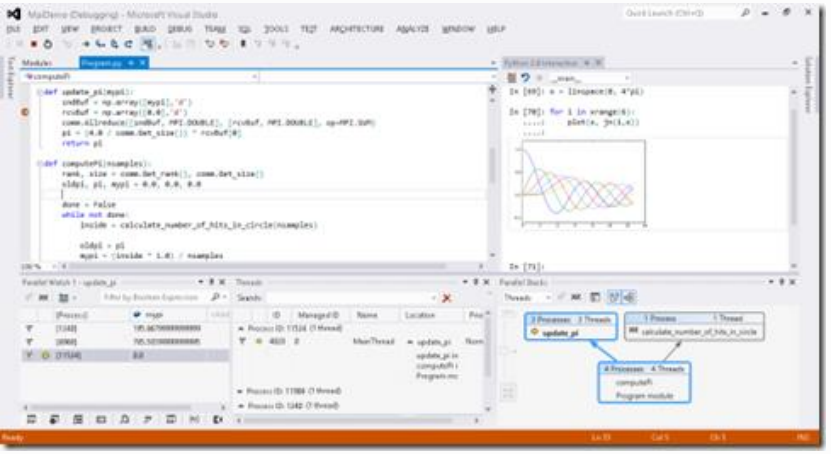

PTVS + Numpy + SciPy is a productive T.C. Workbench

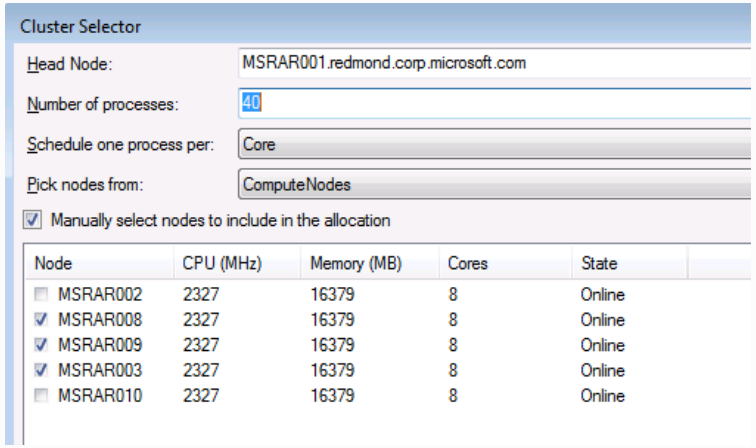

HPC / MPI w support for cluster debugging

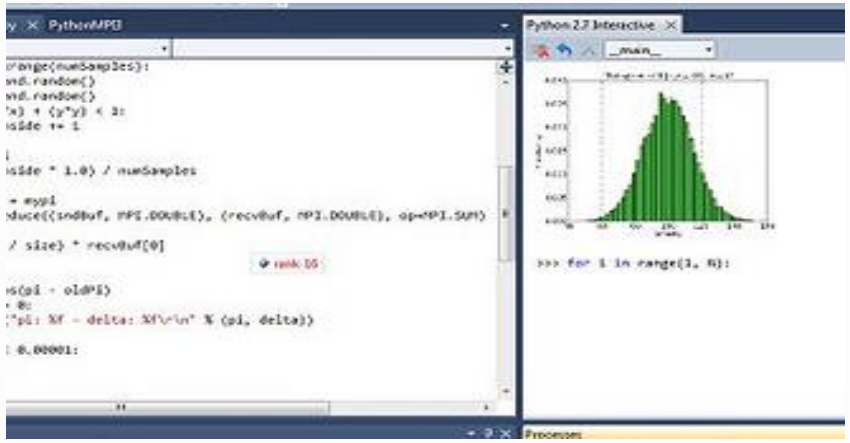

Inline Graphics in REPL

```
In [2]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
        ax1.plot(x); ax1.set_title('Raw audio signal')
        ax2.specgram(x); ax2.set_title('Spectrogram');
```

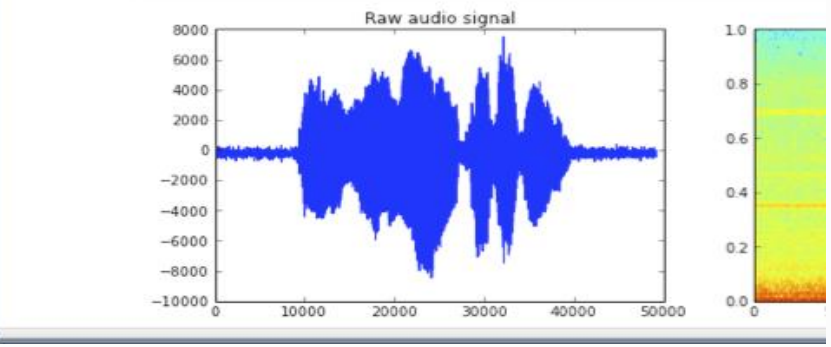

IPython notebook: Python in browser: any OS/ any browser

```
import numpy as np
import numpy.fft as fft

x = np.fromfile('signal.dat', dtype=('>i4,>i2,>i2'))

# Copy, convert to native floats, rescale
sig = np.array(x['f1'], dtype=np.float32) /  32768

# Alternate: since no file is provided, we can make up some data:
data = [math.sin(np.pi / 8 * i) for i in range(320)]
noise = [np.random.rand() * 0.02 for i in range(320)]
sig = np.array(data, dtype='f')
sig += np.array(noise)

f = fft.fft(sig)

print "Strongest frequency = %d n" % np.argmax(np.abs(f))
```

Experimental: .Net versions of Numpy + SciPy

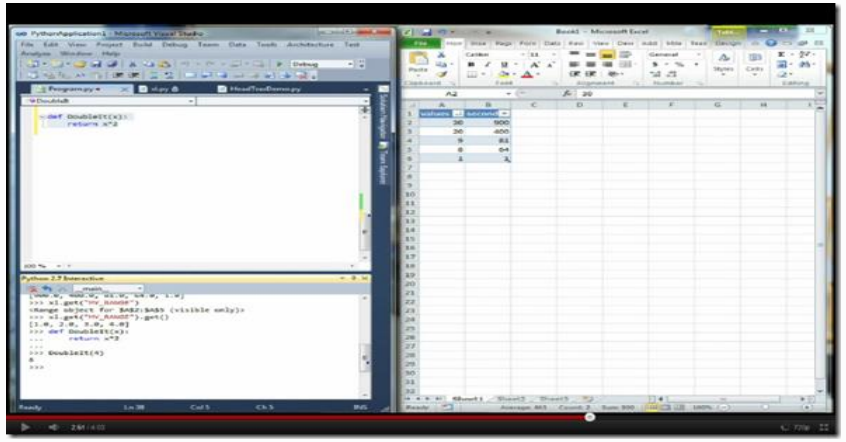

Pyvot: A live bridge between PTVS and Excel

Windows Azure

Python Tools for Visual Studio

PTVS 2.0 Alpha is LIVE! - please install & let us know what you think!

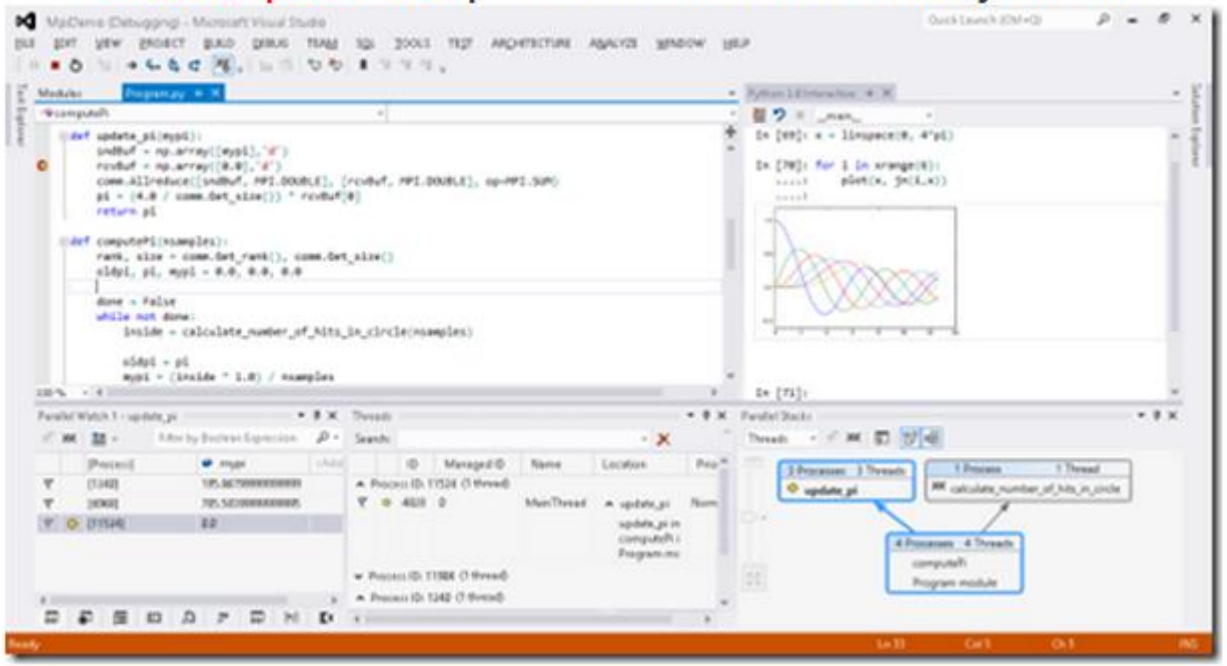

JOIN US

Twitter : @pt4vs http://twitter.com/pt4vs
FB: http://www.facebook.com/pt4vs

Please look at the issue tracker, we're always open to contributions!

Sign in to join this project

Python on Azure: IPython notebook

- Python IDE in browser
  - Any Browser
  - Any OS
- Backed by Python engine on Azure
  - Windows or Linux
- Key features
  - Intellisense, completion, …
  - Inline graphics
  - Markdown
- "Executable Document"
- IPython REPL also built-into PTVS

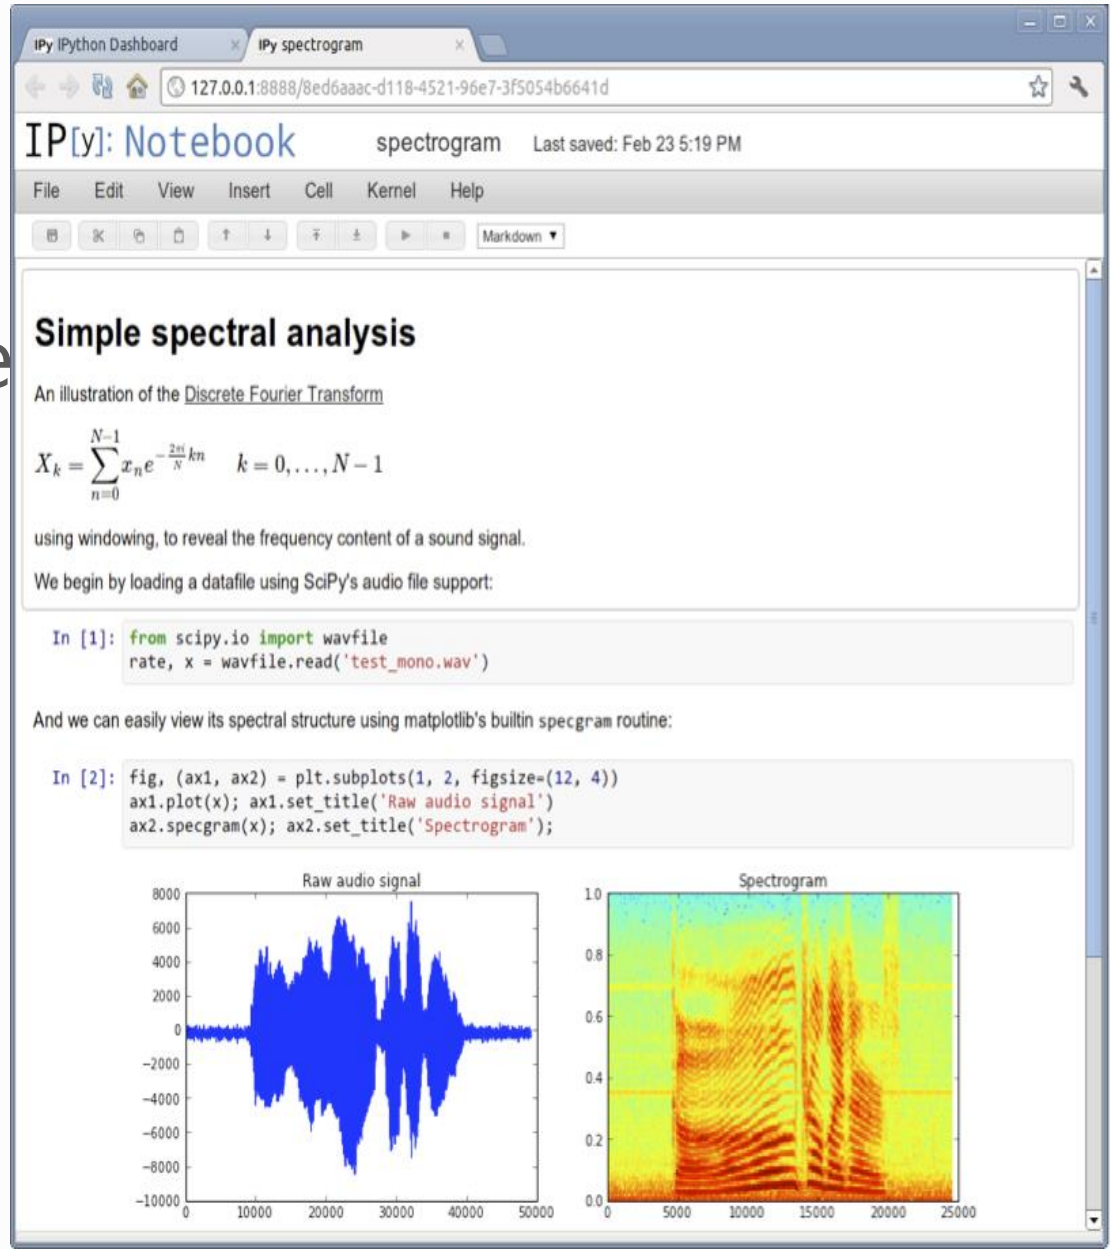

Windows Azure

Python Dashboard

https://dpedemo2.cloudapp.net

P[y]: Notebook

Logout

otebooks | Clusters

import a notebook, drag the file onto the listing below or **click here**.

ome / yew /

demo of K-Means clustering on the handwritten digits data

demo of structured Ward hierarchical clustering on Lena image

troML - Density Estimation

troML - Huber Loss Function

troML - SDSS Stripe 82 Moving Object Catalog

troML - SDSS imaging data and the SDSS Stripe 82 standard stars

Python Dashboard | IPy QuantLib-American Optio

https://dpedemo2.cloudapp.net/65b9e14e-78ab-4775-88c8-c89905b104a3

IP[y]: Notebook

Logo

QuantLib American Option Last saved: Apr 05 5:22 AM

File Edit View Insert Cell Kernel Help

Code Cell Toolbar: None

```
from QuantLib import *

# global data
todaysDate = Date(15,May,1998)
Settings.instance().evaluationDate = todaysDate
settlementDate = Date(17,May,1998)
riskFreeRate = FlatForward(settlementDate, 0.06, Actual365Fixed())

# option parameters
exercise = AmericanExercise(settlementDate, Date(17,May,1999))
payoff = PlainVanillaPayoff(Option.Put, 40.0)

# market data
underlying = SimpleQuote(36.0)
volatility = BlackConstantVol(todaysDate, TARGET(), 0.20, Actual365Fixed())
dividendYield = FlatForward(settlementDate, 0.00, Actual365Fixed())
```

Python Dashboard | IPy QuantLib-American Optio

https://dpedemo2.cloudapp.net/65b9e14e-78ab-4775-88c8-c89905b104a3

P[y]: Notebook

Logo

QuantLib American Option Last saved: Apr 05 5:22 AM

File Edit View Insert Cell Kernel Help

Code Cell Toolbar: None

method	value	estimated error	actual error
reference value	4.48667	n/a	0.0000
Barone-Adesi-Whaley	4.46563	n/a	0.0210
Bjerksund-Stensland	4.45885	n/a	0.0278
finite differences	4.49391	n/a	0.0072
binomial (JR)	4.48655	n/a	0.0001
binomial (CRR)	4.48642	n/a	0.0003
binomial (EQP)	4.48010	n/a	0.0066
bin. (Trigeorgis)	4.48646	n/a	0.0002
binomial (Tian)	4.48641	n/a	0.0003
binomial (LR)	4.48608	n/a	0.0006

IPy[y]: Notebook

QuantLib-American Option

Last saved: Apr 05 5:23 AM


```
process = BlackScholesMertonProcess(QuoteHandle(underlying),
                                    YieldTermStructureHandle(dividendYield),
                                    YieldTermStructureHandle(riskFreeRate),
                                    BlackVolTermStructureHandle(volatility))

option = VanillaOption(payoff, exercise)

refValue = 4.48667344
report('reference value',refValue)

# method: analytic

option.setPricingEngine(BaroneAdesiWhaleyEngine(process))
report('Barone-Adesi-Whaley',option.NPV())

option.setPricingEngine(BjerksundStenslandEngine(process))
report('Bjerksund-Stensland',option.NPV())

# method: finite differences
timeSteps = 801
gridPoints = 800

option.setPricingEngine(FDAmericanEngine(process,timeSteps,gridPoints))
report('finite differences',option.NPV())
```

```
import numpy as np
import pylab as pl
from matplotlib import finance
from matplotlib.collections import LineCollection

from sklearn import cluster, covariance, manifold

###############################################################################
# Retrieve the data from Internet

# Choose a time period reasonnably calm (not too long ago so that we get
# high-tech firms, and before the 2008 crash)
d1 = datetime.datetime(2003, 01, 01)
d2 = datetime.datetime(2008, 01, 01)

symbol_dict = {
    'TOT': 'Total',
    'XOM': 'Exxon',
    'CVX': 'Chevron',
    'COP': 'ConocoPhillips',
    'VLO': 'Valero Energy',
    'MSFT': 'Microsoft',
    'IBM': 'IBM',
```

Python Dashboard | IPy Visualizing the stock mark | IPy QuantLib-American Optic

https://dpedemo2.cloudapp.net/3515c455-c06f-4b68-b01e-78a0da7ab8d0

[y]: Notebook

Visualizing the stock market structure

Last saved: Apr 05 5:25 AM

Edit View Insert Cell Kernel Help

Code | Cell Toolbar: None

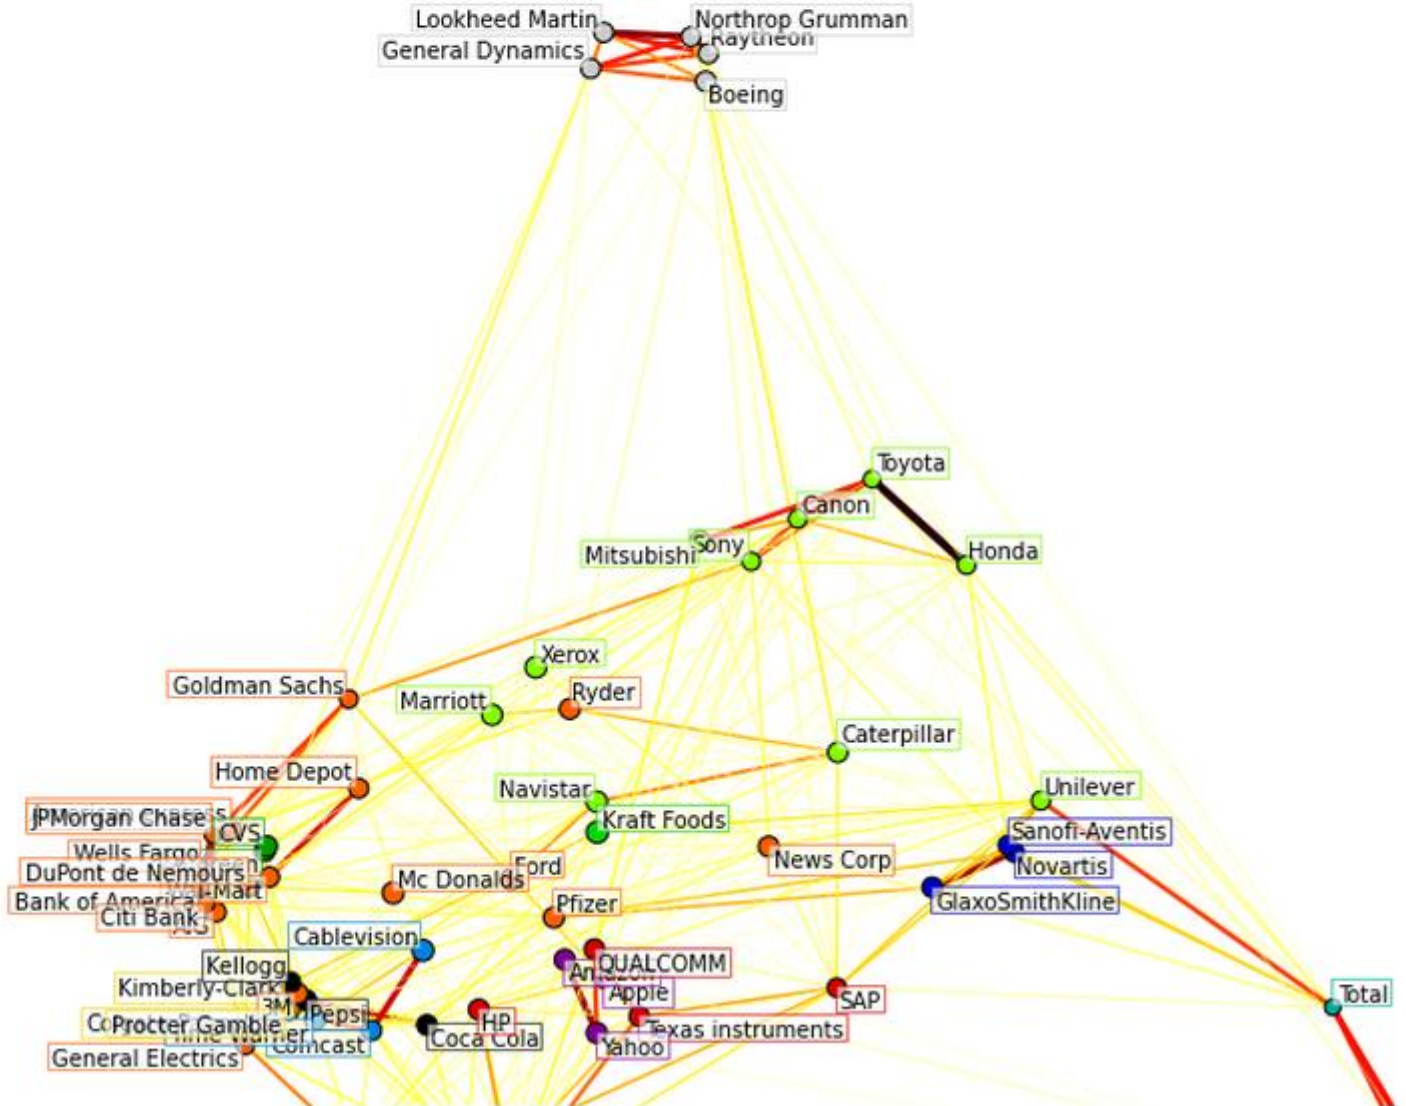

Virtual machines

Windows Server and Linux

Flexible Workload Support

Virtual Private Networking

Virtual machine operating system selection

- ALL
- PLATFORM IMAGES
- MY IMAGES
- MY DISKS

- Windows Server 2008 R2 SP1
- Windows Server 2012 Datacenter
- OpenLogic CentOS 6.3
- openSUSE 12.3
- RightScale Linux v13
- SUSE Linux Enterprise Server 11 SP2
- Ubuntu 12.04
- Ubuntu 12.10
- wenmingsaved
- whitehall
- boothdemo1-boothdemo1-0-2012

Microsoft SQL Server...

SQL Server 2012 SP1 Cumulative Update 2 Evaluation Edition (64-bit) on Windows Server 2008 R2 Service Pack 1. Virtual Machines created by using this SQL Server Evaluation Edition will expire on August 20, 2013. This image contains the full version of SQL Server. Some SQL Server 2012 components require additional setup and configuration before use. Medium is the minimum recommended virtual machine size for this image. To evaluate the advanced capabilities of SQL Server 2012, we recommend that you use a virtual machine size of Large or Extra Large.

PUBLISHER	Microsoft SQL Server Group
OS FAMILY	Windows
LOCATION	East Asia;Southeast Asia;North Europe;West Europe;East US;West US

2 3 4

CREATE A VIRTUAL MACHINE

Virtual machine configuration

VERSION RELEASE DATE

February 27, 2013

VIRTUAL MACHINE NAME

NEW USER NAME

azureuser

NEW PASSWORD

CONFIRM

SIZE

Extra Large (8 cores, 14 GB Memory)

UPLOAD SSH KEY FOR AUTHENTICATION

Ubuntu 12.10

Ubuntu Server 12.10 (amd64 20130227) for Windows Azure. Ubuntu Server is the world's most popular Linux for cloud environments. Updates and patches for Ubuntu 12.10 will be available until April 2014. Ubuntu Server is the perfect platform for all workloads from web applications to NoSQL databases and Hadoop. More information can be found at: http://www.ubuntu.com/business/server

PUBLISHER	Canonical
OS FAMILY	Linux
LOCATION	East Asia;Southeast Asia;North Europe;West Europe;East US;West US

WindowsAzure / azure-sdk-for-python

Unwatch | Star 60 | For

Code | Network | Pull Requests 0 | Issues 13 | Wiki | Graphs

Windows Azure SDK for Python — Read more

Clone in Windows | ZIP | HTTP | SSH | Git Read-Only | git@github.com:WindowsAzure/azure-sdk-for-pytl | Read-Only

branch: master | Files | Commits | Branches 2

zure-sdk-for-python /

43 c

Merge pull request #77 from WindowsAzure/dev ...

jeffwilcoxmsft authored 15 days ago — latest commit 738c79

src	16 days ago	Updated version for next release [huguesv]
test	15 days ago	Split normal run and coverage run in separate batch files and added a... [huguesv]

Deploy & run Python on Virtual Machine

Demo

```
from azure import *
from azure.servicemanagement import *

sms = ServiceManagementService(subscription_id, certificate_path)

name = 'myvm'
location = 'West US'

# You can either set the location or an affinity_group
sms.create_hosted_service(service_name=name, label=name, location=location)

# Name of an os image as returned by list_os_images
image_name = 'OpenLogic__OpenLogic-CentOS-62-20120531-en-us-30GB.vhd'

# Destination storage account container/blob where the VM disk
# will be created
media_link = 'url_to_target_storage_blob_for_vm_hd'

# Linux VM configuration, you can use WindowsConfigurationSet
# for a Windows VM instead
linux_config = LinuxConfigurationSet('myhostname', 'myuser', 'mypassword', True)

os_hd = OSVirtualHardDisk(image_name, media_link)

sms.create_virtual_machine_deployment(service_name=name,
    deployment_name=name, deployment_slot='production', label=name,
    role_name=name, system_config=linux_config, os_virtual_hard_disk=os_hd, role_size='Small')
```

Virtual machines

Gallery Images

Virtual Machine Portability

VMs with persistent Drives

Gallery Images Available

Microsoft

Windows Server 2008 R2

SQL Server Eval 2012

Windows Server 2012

Biztalk Server 2013 Beta

Microsoft® SQL Server® 2012

Microsoft® BizTalk® Server

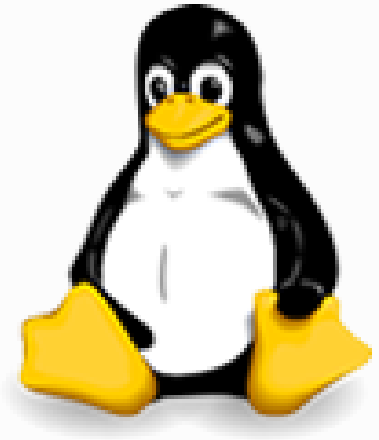

Open Source

OpenSUSE 12.2

CentOS 6.3

Ubuntu 12.04/12.10

SUSE Linux Enterprise Server 11 SP2

Windows Azure™

Virtual machine portability

Windows Azure

Your Data Center

Windows Azure

Windows Azure

Your Data Center

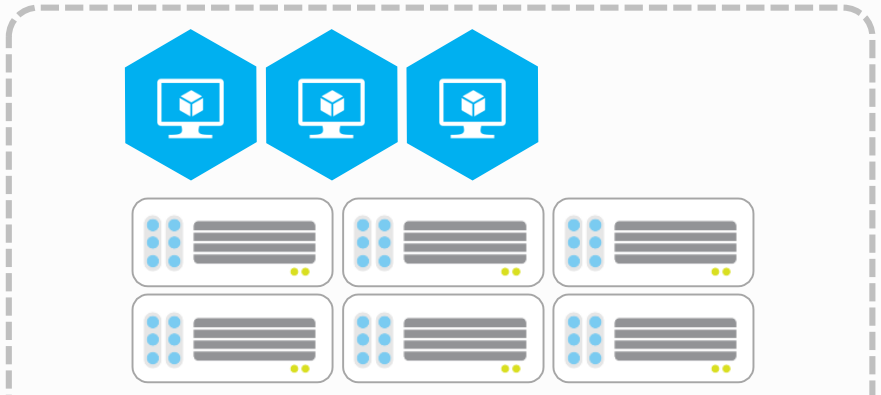

Windows Azure

Windows Azure

Other Service Providers

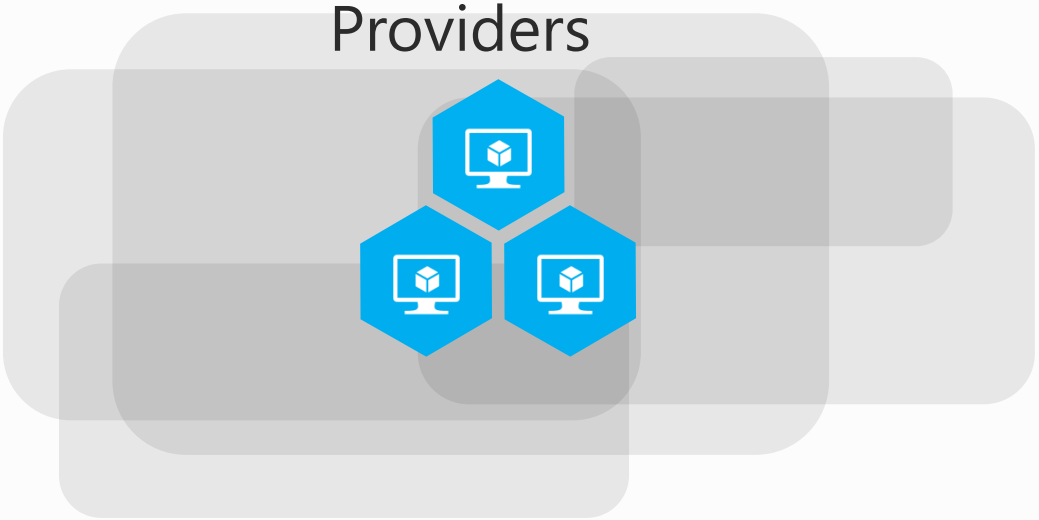

Your Data Center

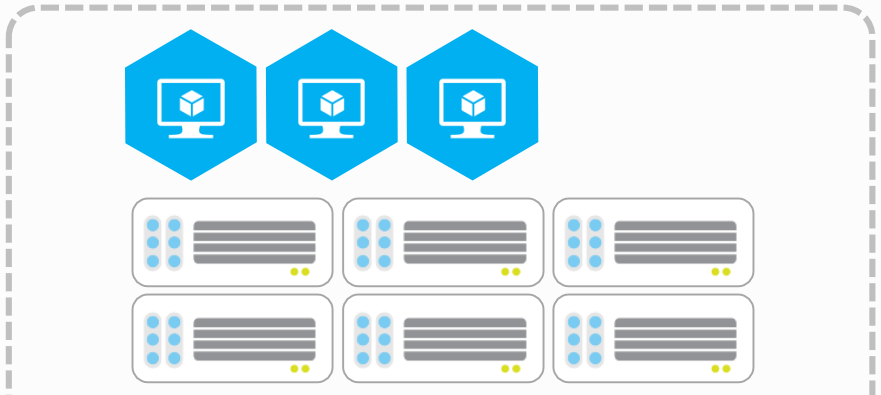

Windows Azure

VM with persistent drive

Windows Azure Storage

VM with persistent drive

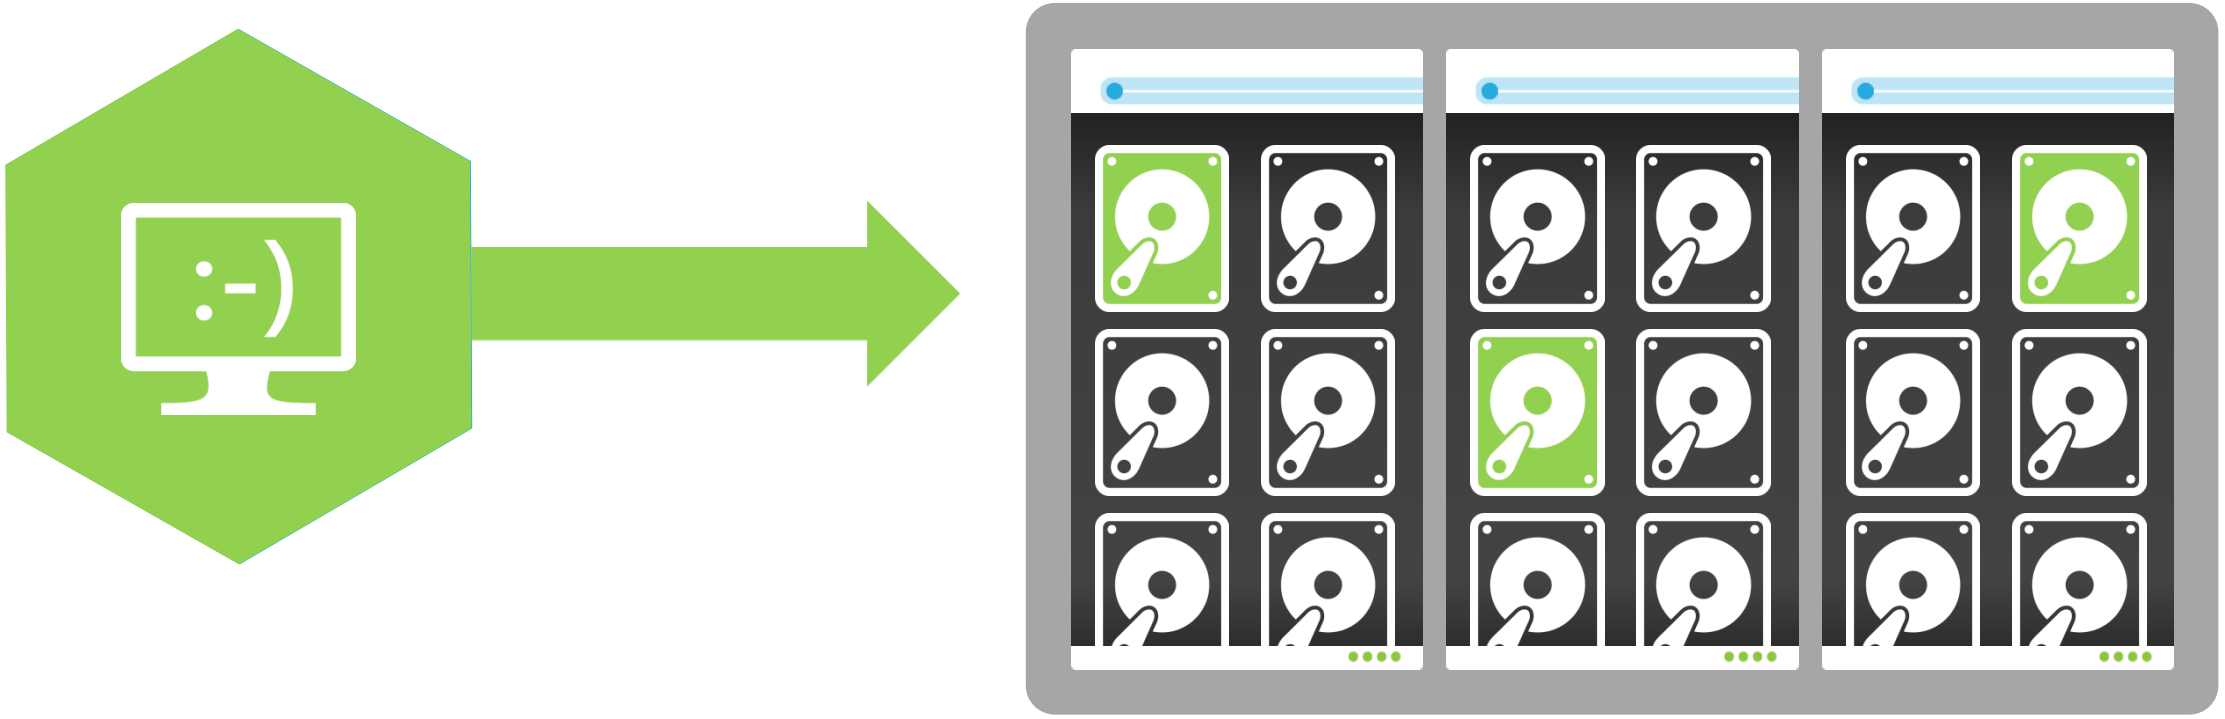

VM with persistent drive

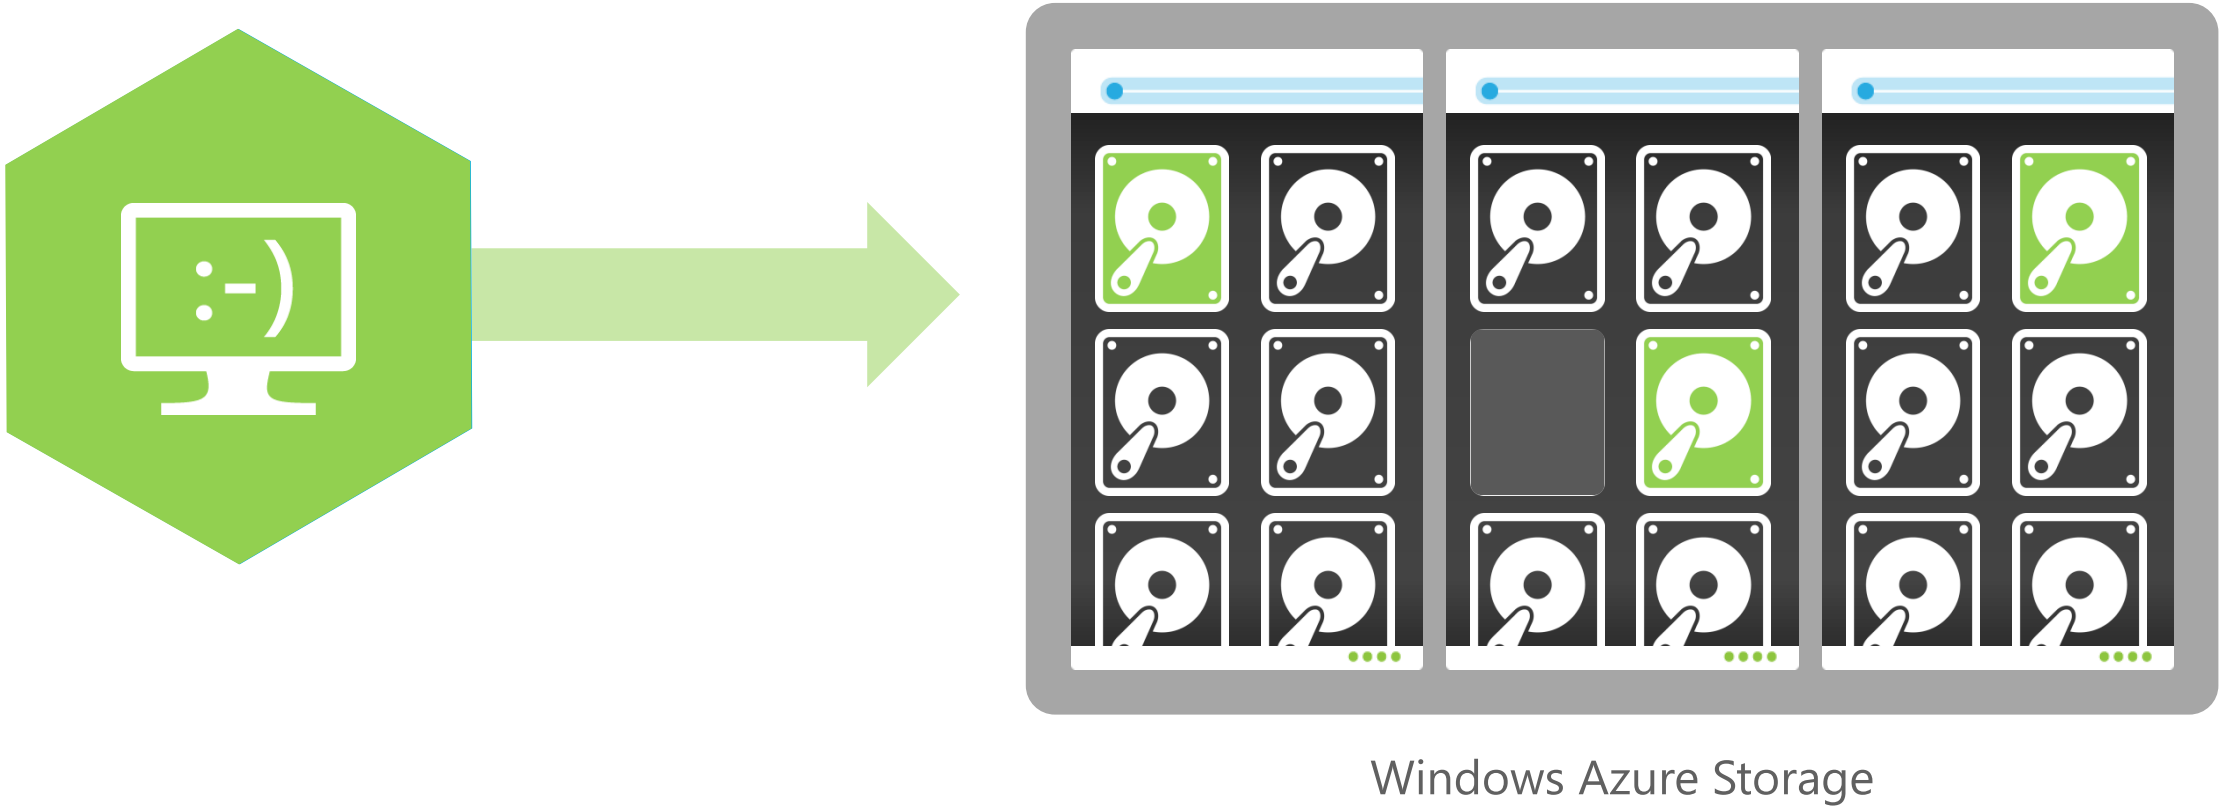

Cloud services

Build infinitely scalable apps and services

Support rich multi-tier architectures

Automated application management

Cloud services: Development

Demo

http://weatherservice.cloudapp.net/Gallery/List

List - Weather Forecasting ...

Storage - Windows Azure ... | Suggested Sites | Concur Hotel Reservation... | Windows Azure HDInsight

NORTHWIND WEATHER RESEARCH

HOME / GALLERY / NEW FORECAST / PROCESSING QUEUE

LIST | MAP | TIMELINE

Sydney Harbour

4/5/2013 12:00:00 AM

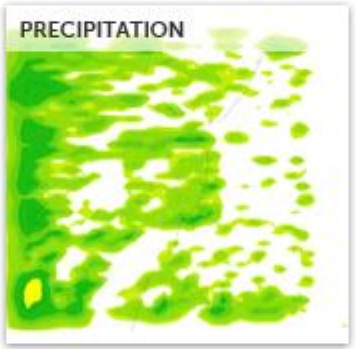

ny

4/4/2013 12:00:00 PM

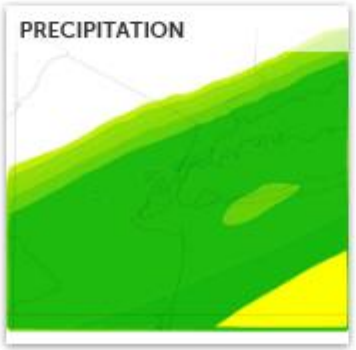

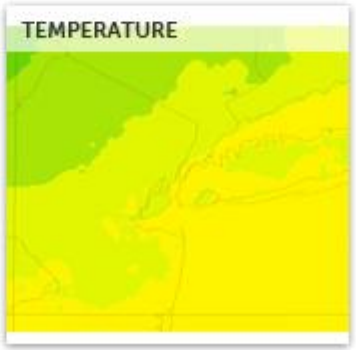

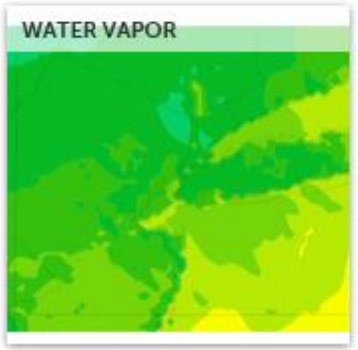

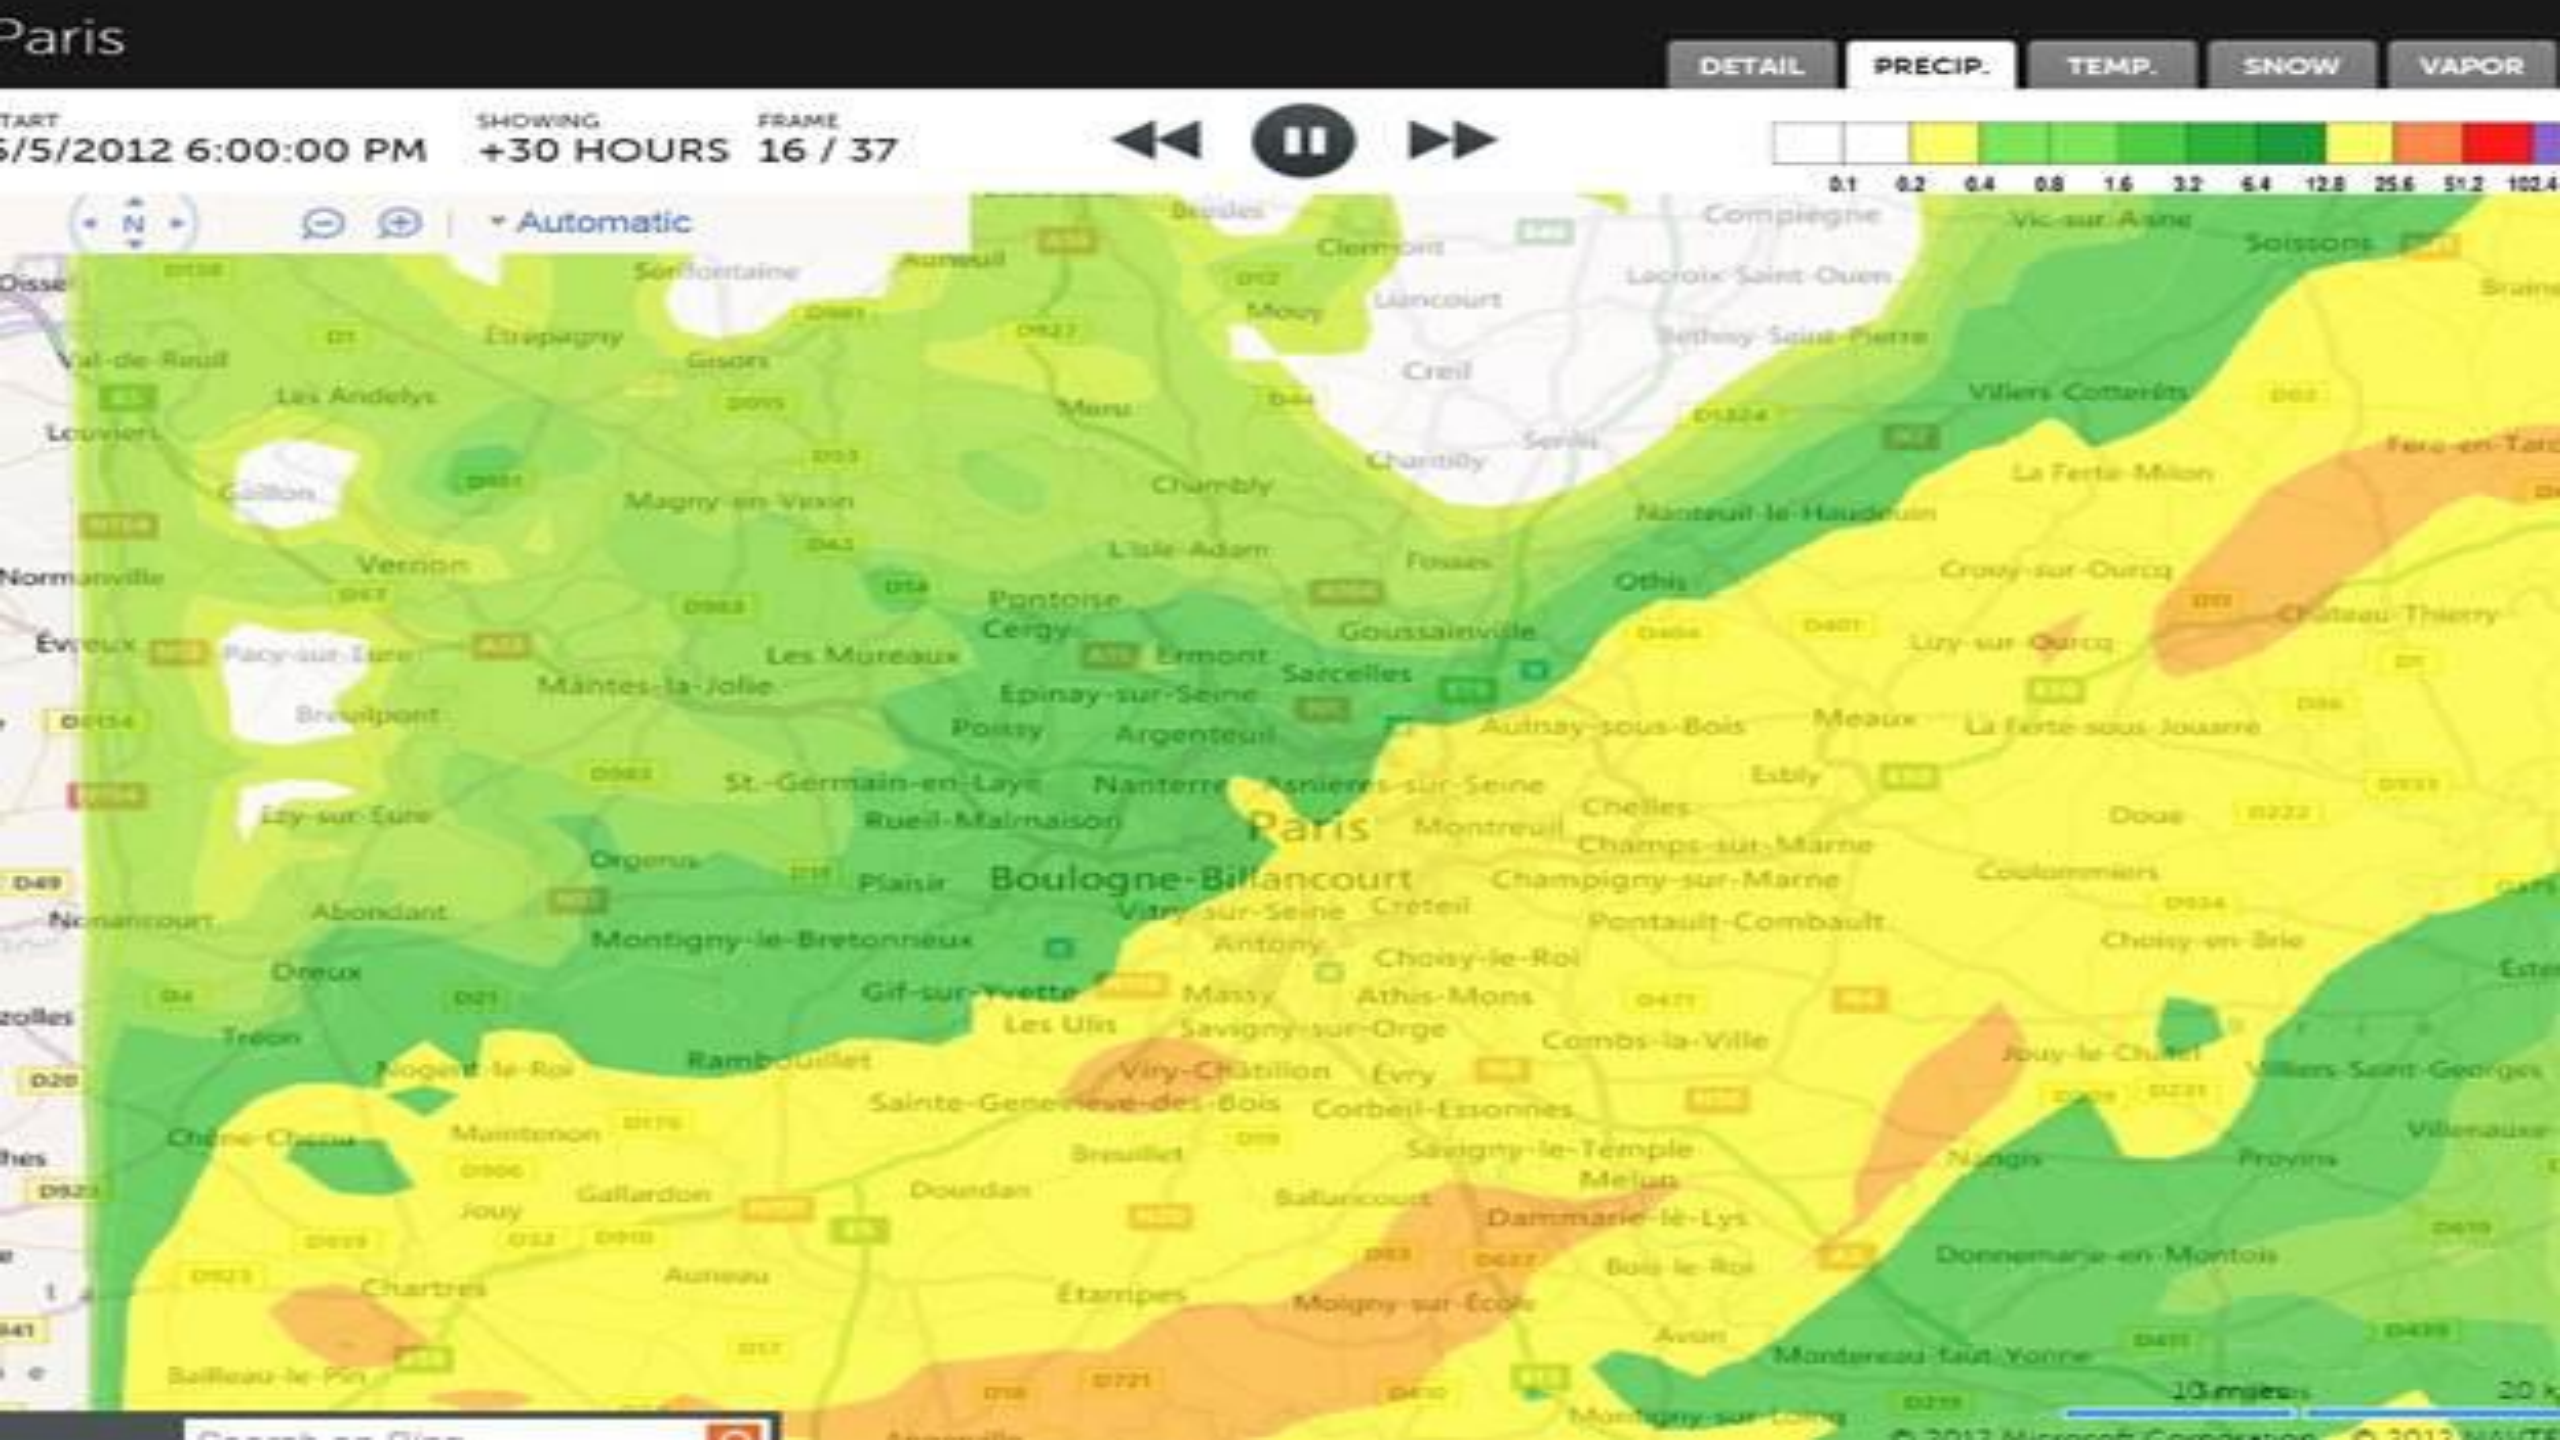
Paris
DETAIL
PRECIP.
TEMP.
SNOW
VAPOR
START
5/5/2012 6:00:00 PM
SHOWING
+30 HOURS
FRAME
16 / 37
0.1
0.2
0.4
0.8
1.6
3.2
6.4
12.8
25.6
51.2
102.4
Automatic
Paris
Boulogne-Billancourt
Versailles
Saint-Germain-en-Laye
Nanterre
Argenteuil
Pontoise
Cergy
Meaux
Melun
Évreux
Dreux
Chartres
Étampes
Fontainebleau
Montereau-fault-Yonne
10 miles

weatherservice

DASHBOARD MONITOR CONFIGURE **SCALE** INSTANCES LINKED RESOURCES CERTIFICATES

PRODUCTION STAGING

roles

HEADNODE
MEDIUM VM

1 instance(s)

FRONTEND2
MEDIUM VM

1 instance(s)

COMPUTENODE
EXTRA LARGE VM

12 instance(s)

SAVE DISCARD

```
# Create container
from azure.storage import BlobService
blob_service = BlobService(account_name, account_key)
blob_service.create_container('taskcontainer')

# Upload
from azure.storage import BlobService
blob_service = BlobService(account_name, account_key)
blob_service.put_blob('taskcontainer', 'task1',
file('task1-upload.txt').read(), 'BlockBlob')

#Download
from azure.storage import BlobService
blob_service = BlobService(account_name, account_key)
blob = blob_service.get_blob('taskcontainer', 'task1')
```

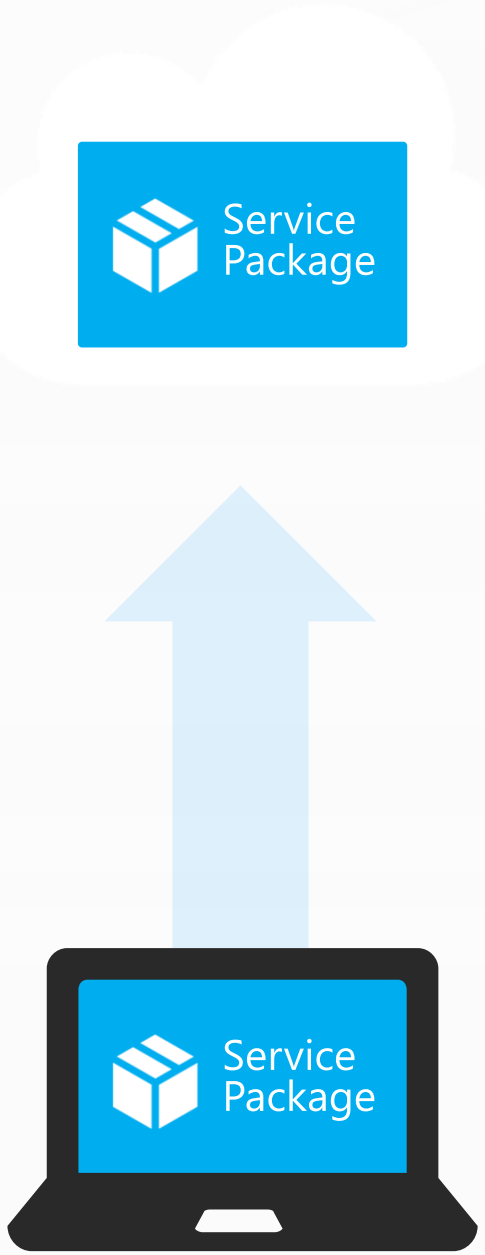
Windows Azure
Service Package
Service Package

Provision Role Instances

Deploy App Code

Configure Network

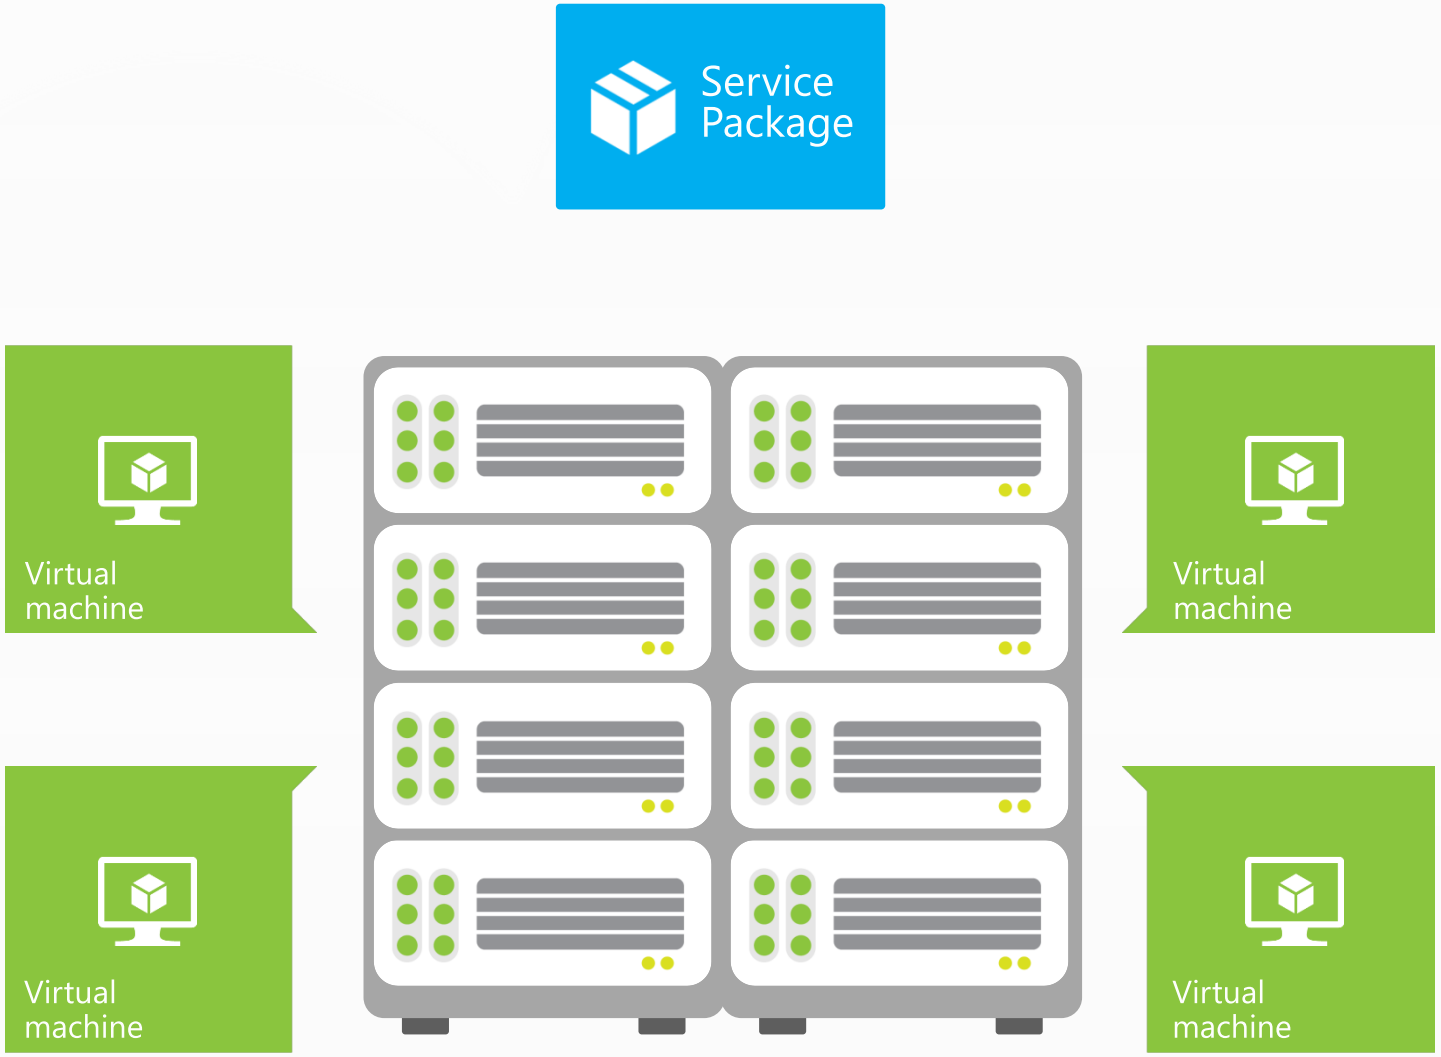

Windows Azure

Provision Role Instances

Deploy App Code

Configure Network

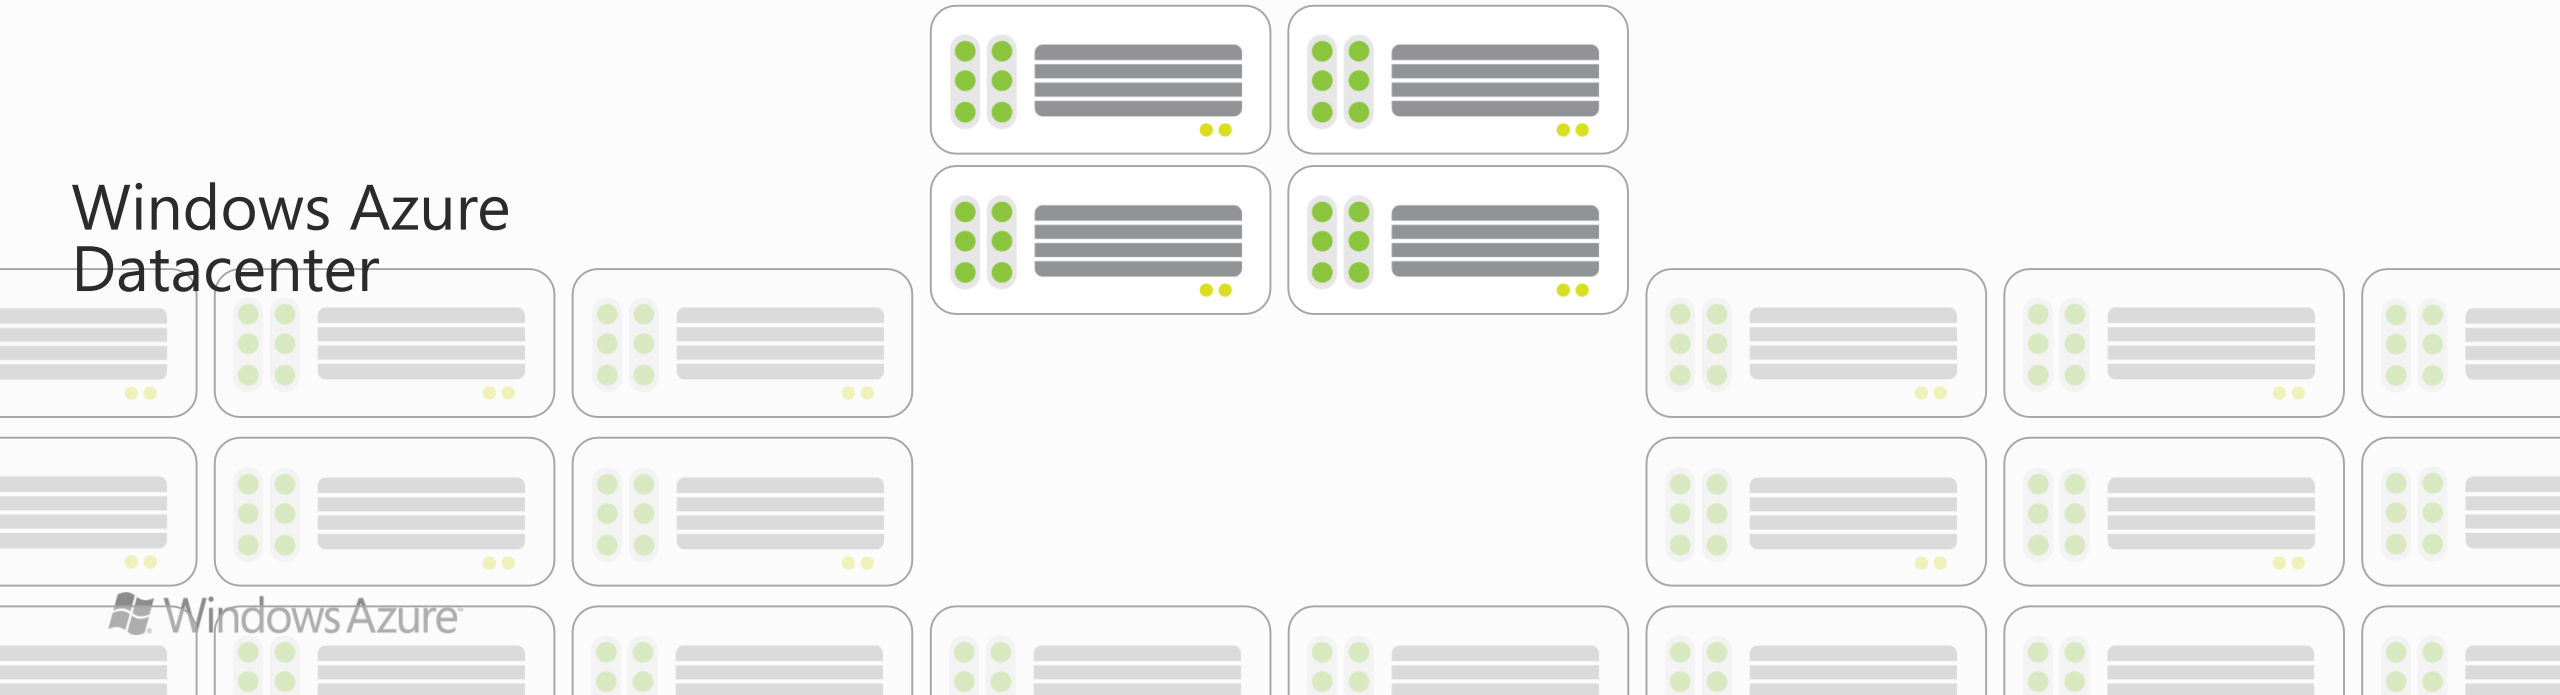

Provision Role Instances

Deploy App Code

Configure Network

Windows Azure
Datacenter

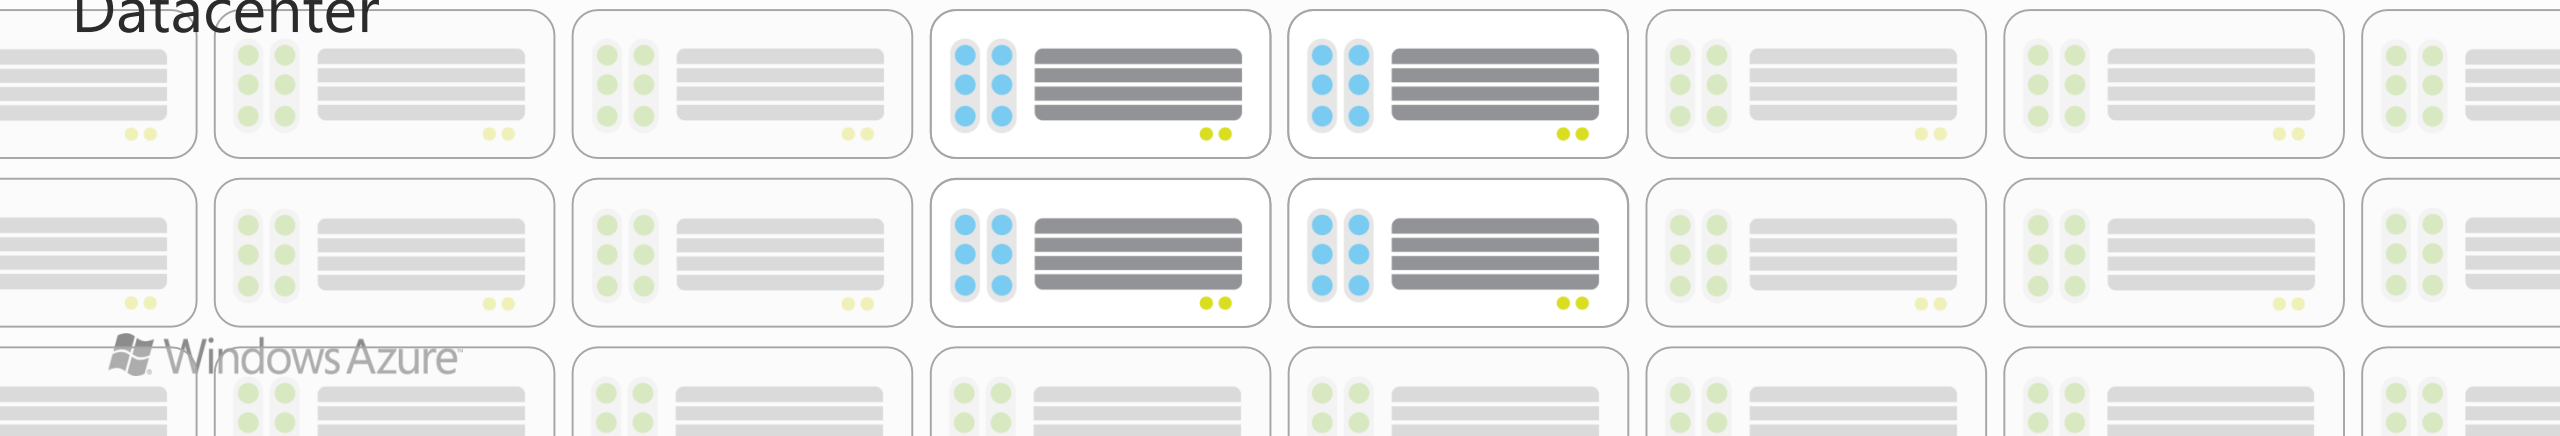

Windows Azure

Provision Role Instances
Deploy App Code
Configure Network

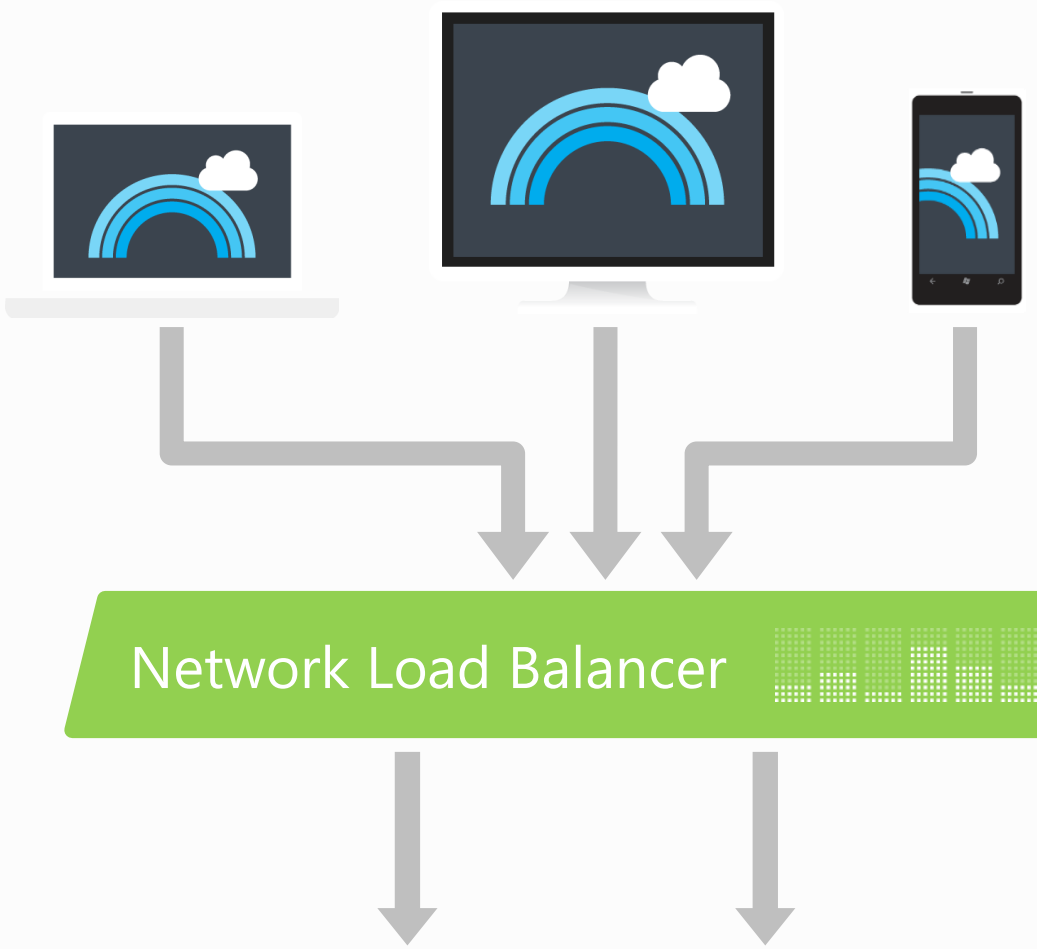

← Network load-balancer configured for traffic

Windows Azure Datacenter

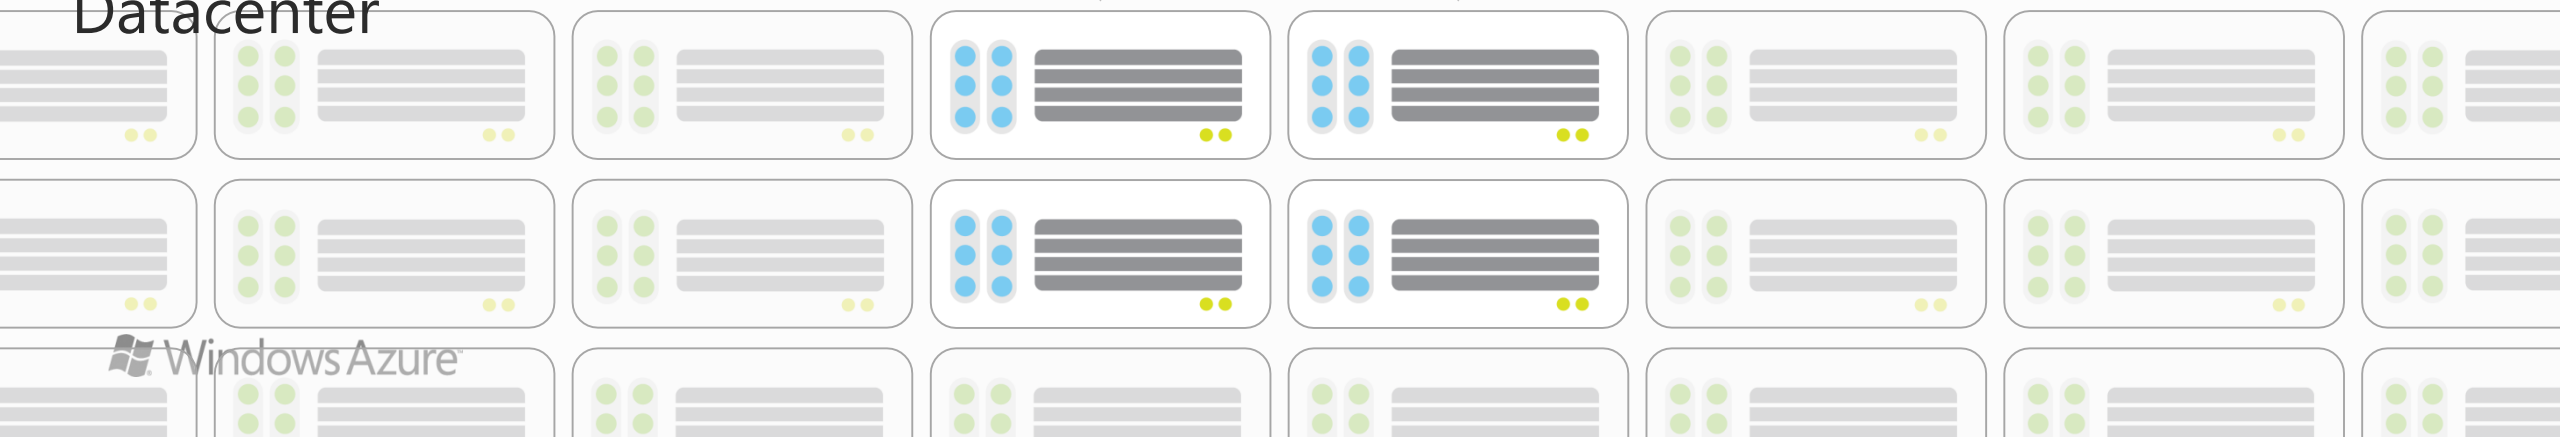

Windows Azure

Application building blocks

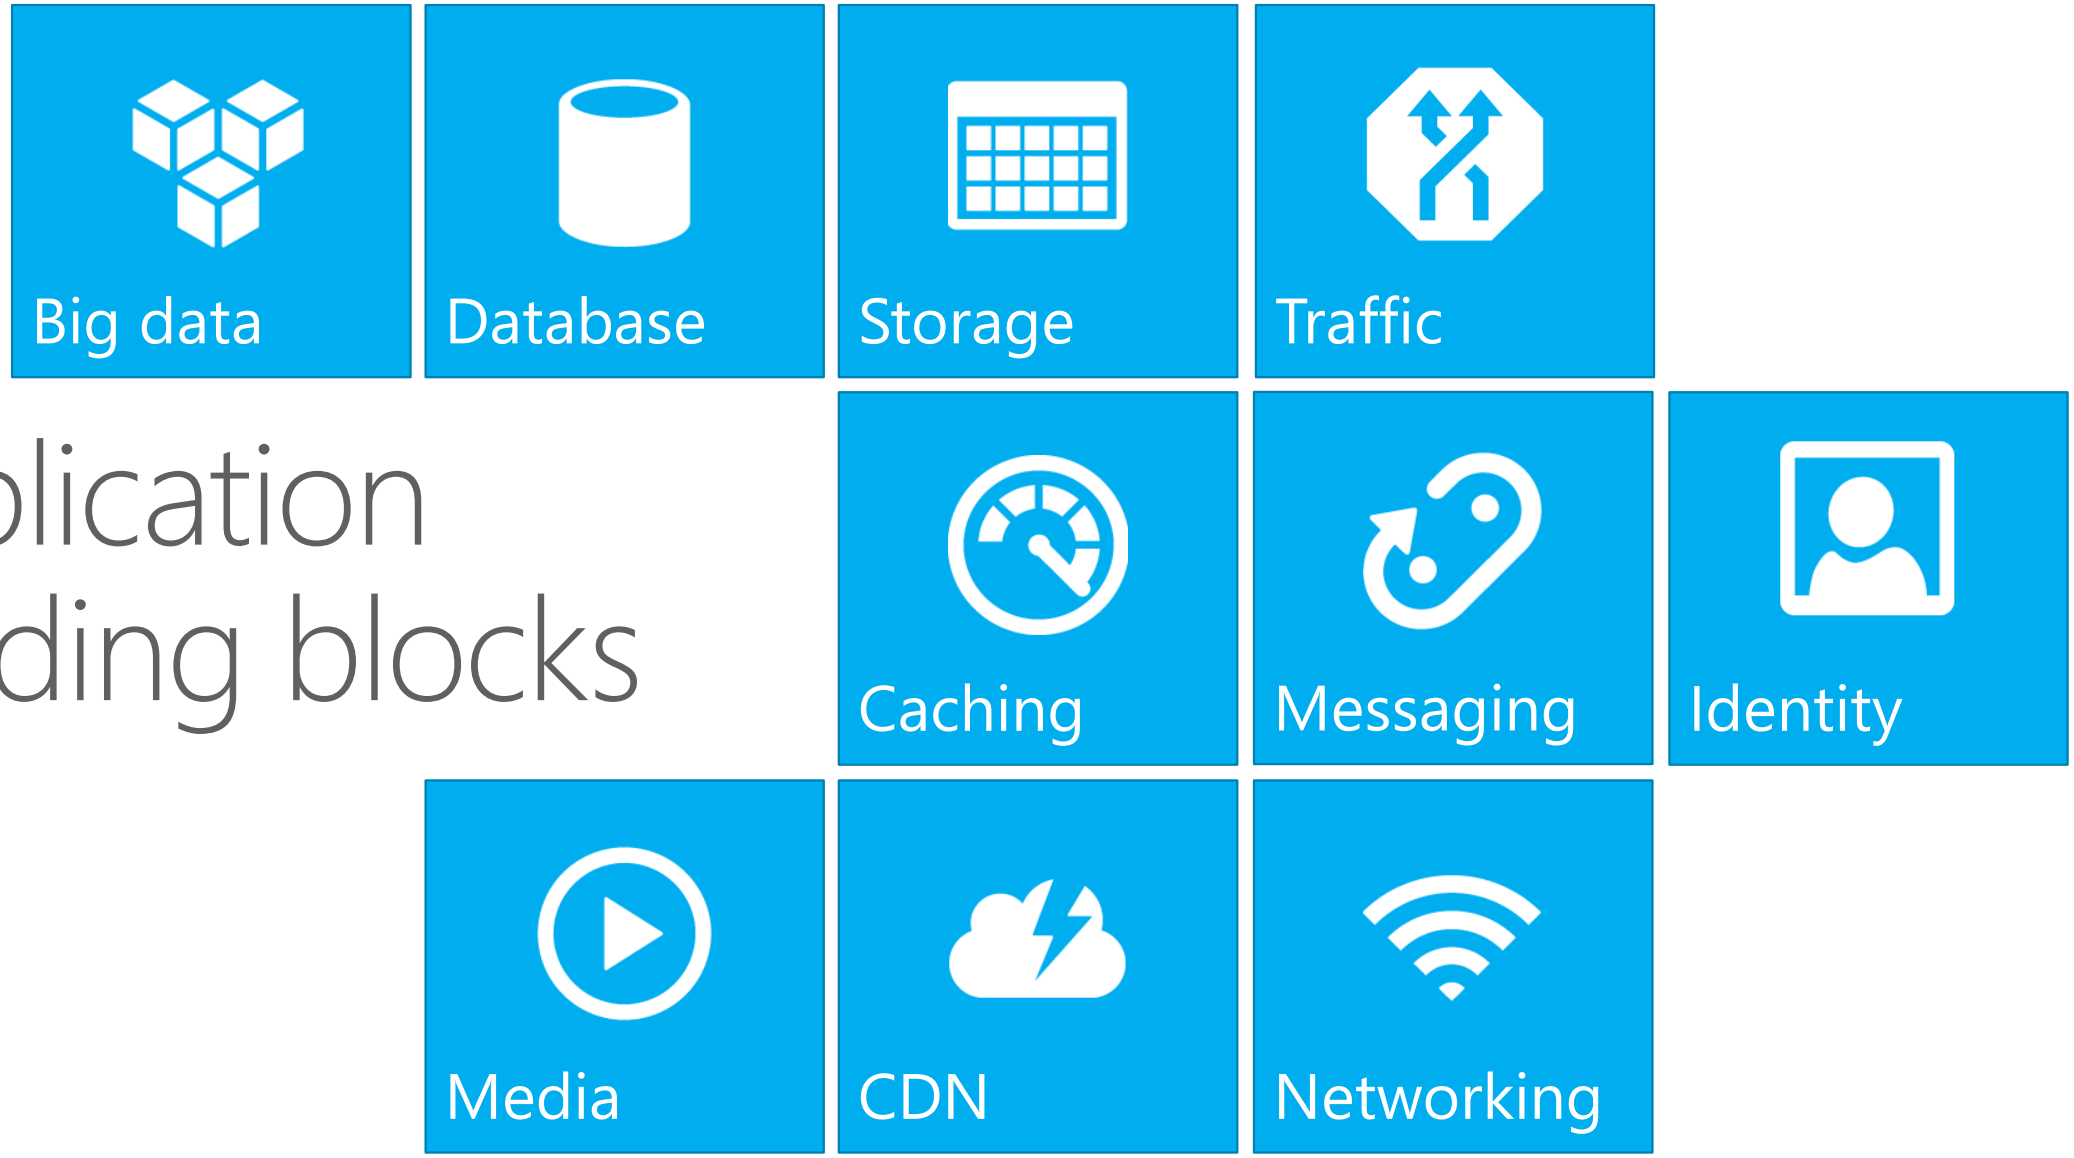

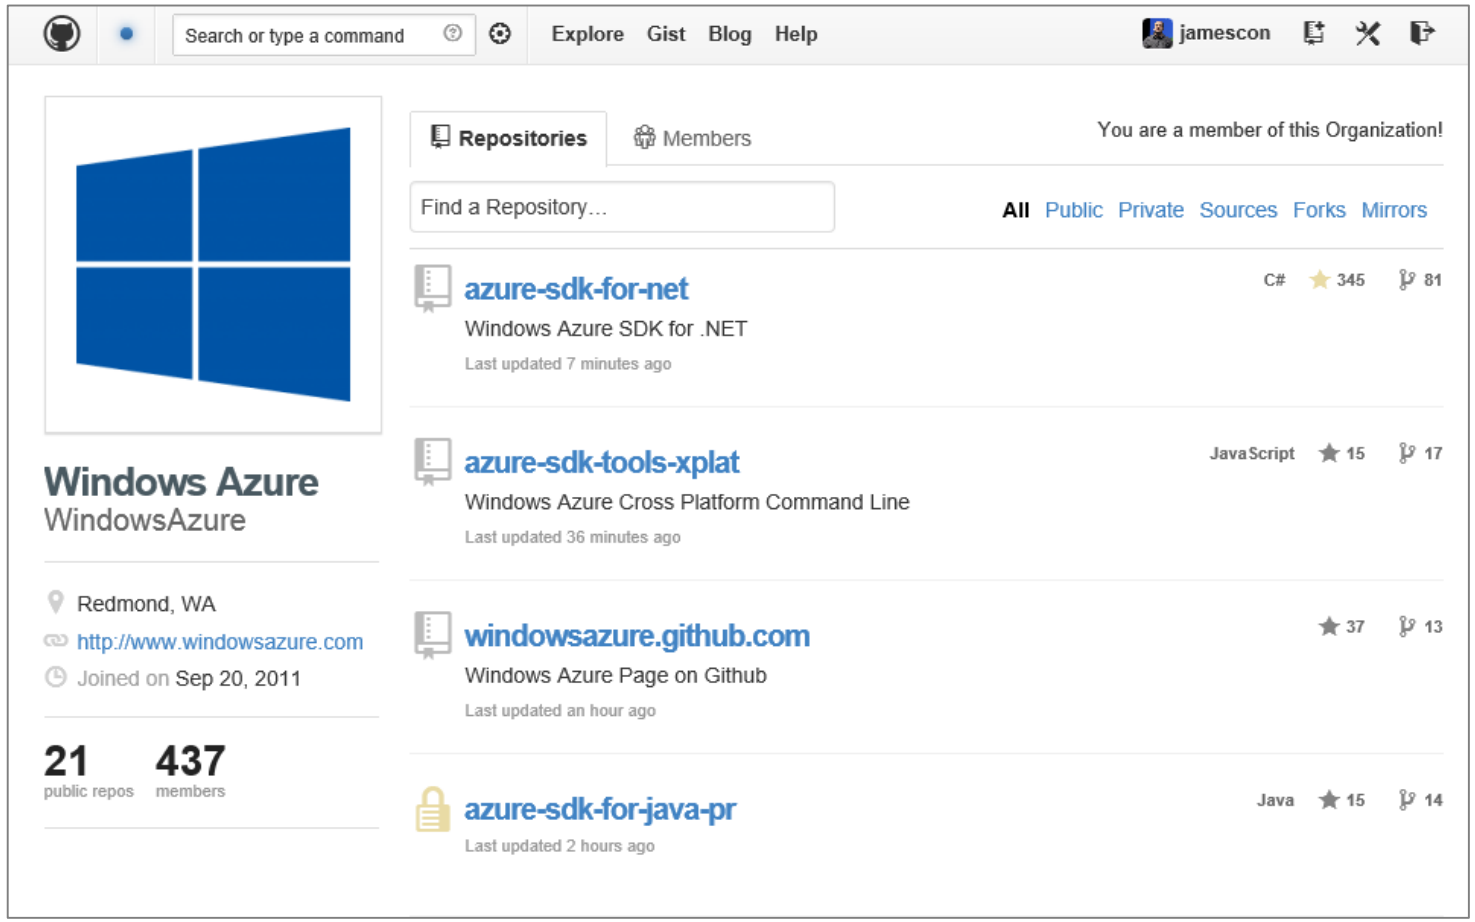

Open source

→ http://github.com/windowsazure

Windows Azure Scenarios

Ideal for Applications Needing:

Scalability
Availability
Fault Tolerance

Common Application Uses:

Web Sites
Compute Intensive apps
Device Applications
Web APIs
Social Games

89 countries and territories

Australia
Austria
Belgium
Brazil
Canada
Chile
Colombia
Costa Rica
Cyprus
Czech Republic
Denmark
Finland
France
Germany
Greece
Hong Kong
Hungary
India
Ireland
Israel
Italy
Japan
Korea
Luxembourg
Malaysia
Mexico
Netherlands
New Zealand
Norway
Peru
Philippines
Poland
Portugal
Puerto Rico
Romania
Russia
Singapore
Spain
Sweden
Switzerland
Trinidad & Tobago
UK
United States

New Countries:

Algeria
Argentina
Belarus
Bulgaria
Croatia
Dominican Rep
Ecuador
Egypt
El Salvador
Estonia
Guatemala
Iceland
Indonesia
Jordan
Kazakhstan
Kenya
Kuwait
Latvia
Liechtenstein
Lithuania
Macedonia
Malta
Montenegro
Morocco
Azerbaijan
Nigeria
Oman
Pakistan
Panama
Paraguay
Qatar
Saudi Arabia
Serbia
Slovakia
Slovenia
South Africa
Sri Lanka
Taiwan
Thailand
Tunisia
Turkey
UAE
Ukraine
Uruguay
Venezuela
Bahrain

Windows Azure

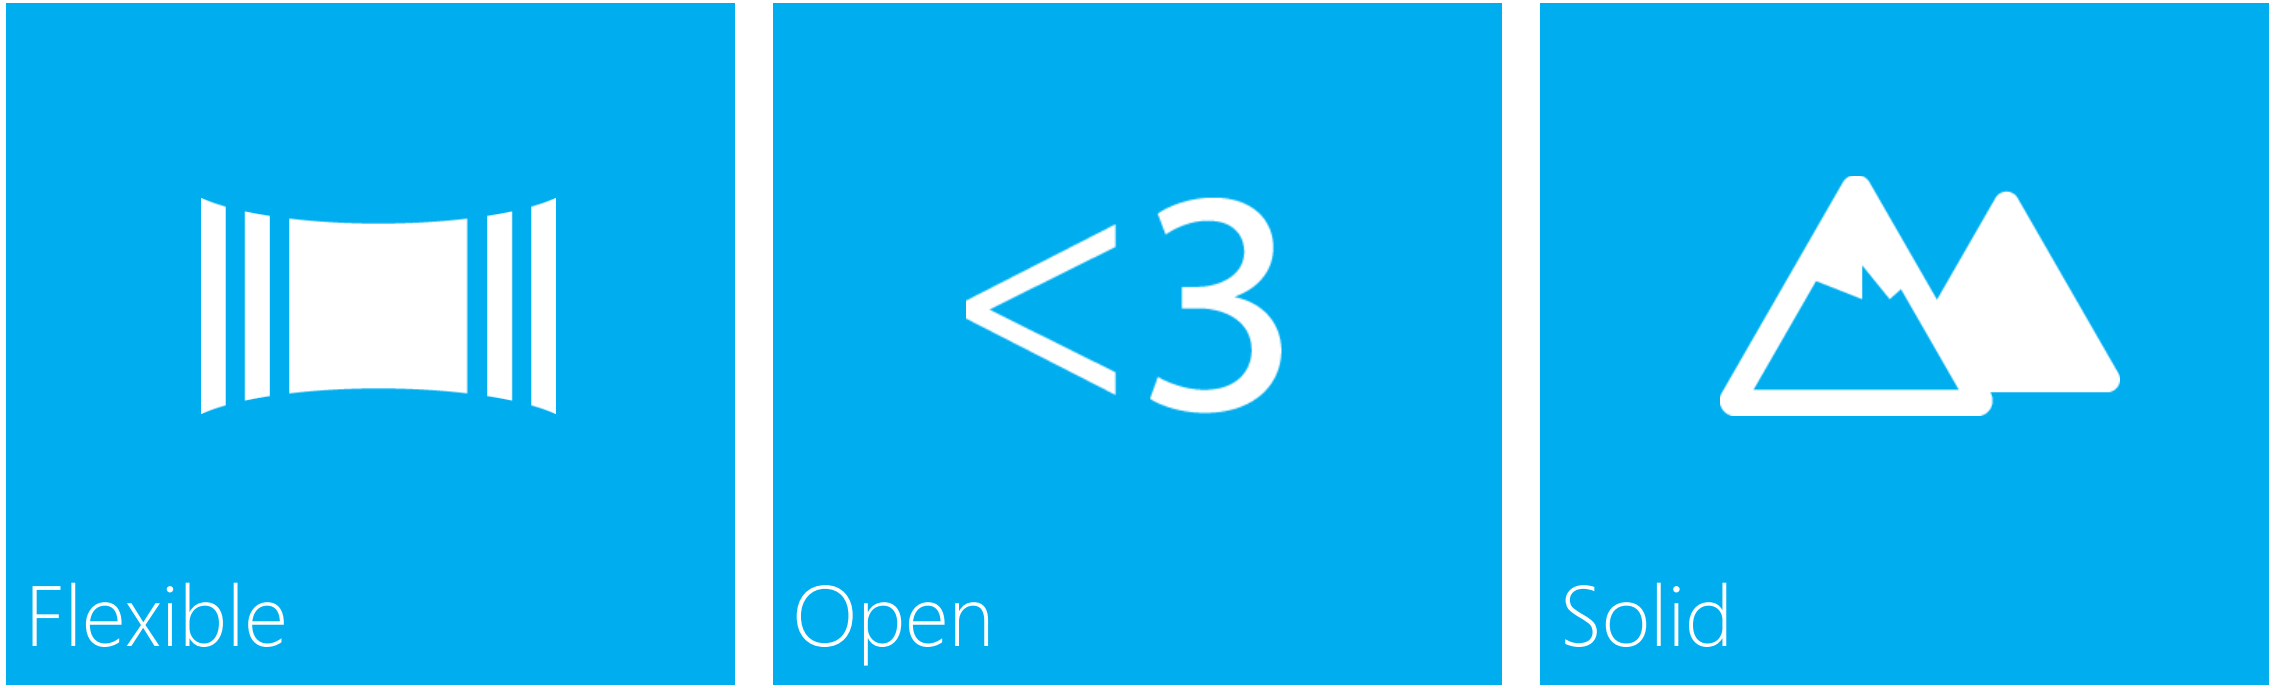

Summary

Windows Azure provides a comprehensive set of services that you can selectively compose to build your cloud apps

Global Data Center Footprint

99.95% Monthly SLA. Pay only for what you use.

Flexible & Open Compute Options

Virtual Machines, Web Sites, & Cloud Services

Managed Building Block Services

SQL Database, Cache, Service Bus, & more

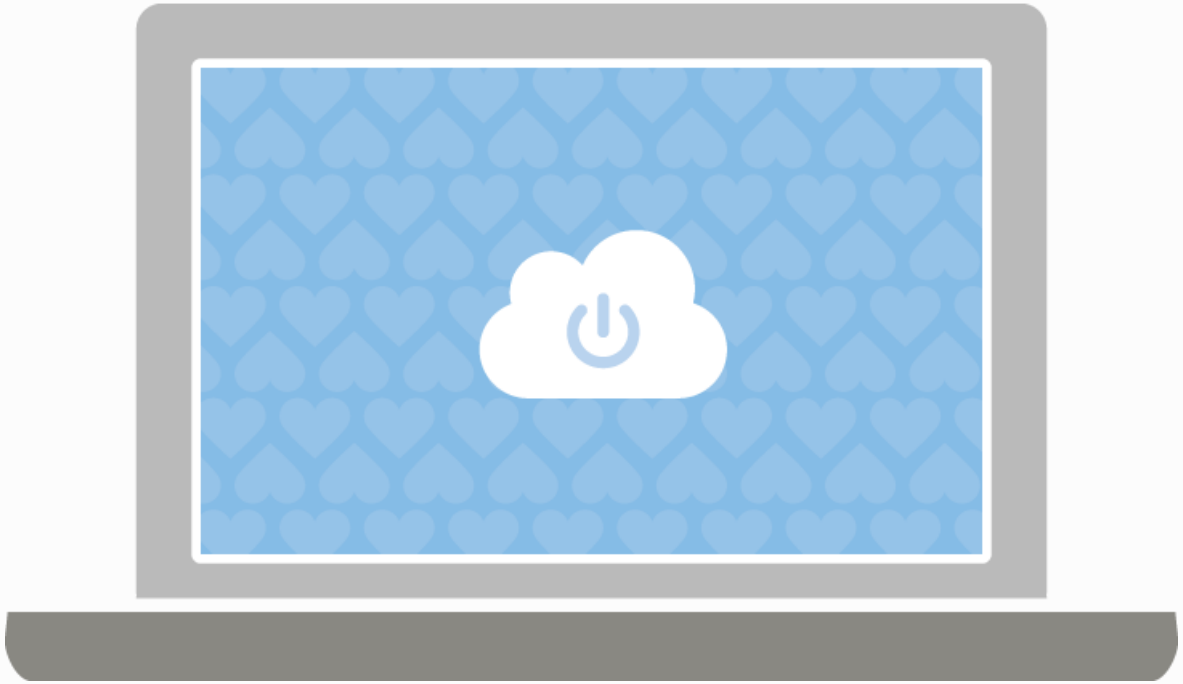

References

- http://pytools.codeplex.com
- Windows Azure Python SDK
- Windows Azure
- How to use Service Management from Python
- http://research.microsoft.com/en-us/projects/azure/

Microsoft®

© 2011 Microsoft Corporation. All rights reserved. Microsoft, Windows, Windows Vista and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries.
The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.